

# Personalized cancer diagnosis

## 1. Business Problem

### 1.1. Description

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/>

Data: Memorial Sloan Kettering Cancer Center (MSKCC)

Download training\_variants.zip and training\_text.zip from Kaggle.

**Context:**

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/discussion/35336#198462>

**Problem statement :**

Classify the given genetic variations/mutations based on evidence from text-based clinical literature.

### 1.2. Source/Useful Links

Some articles and reference blogs about the problem statement

1. <https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>  
(<https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>)
2. <https://www.youtube.com/watch?v=UwbuW7oK8rk> (<https://www.youtube.com/watch?v=UwbuW7oK8rk>)
3. <https://www.youtube.com/watch?v=qxXRKVompl8> (<https://www.youtube.com/watch?v=qxXRKVompl8>)

### 1.3. Real-world/Business objectives and constraints.

- No low-latency requirement.
- Interpretability is important.
- Errors can be very costly.
- Probability of a data-point belonging to each class is needed.

## 2. Machine Learning Problem Formulation

### 2.1. Data

#### 2.1.1. Data Overview

- Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>  
(<https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>)
- We have two data files: one contains the information about the genetic mutations and the other contains the clinical evidence (text) that human experts/pathologists use to classify the genetic mutations.
- Both these data files have a common column called ID
- Data file's information:
  - training\_variants (ID , Gene, Variations, Class)
  - training\_text (ID, Text)

#### 2.1.2. Example Data Point

**training\_variants**

ID, Gene, Variation, Class

0, FAM58A, Truncating Mutations, 1

1, CBL, W802\*, 2

2, CBL, Q249E, 2

...

**training\_text**

ID, Text

0||Cyclin-dependent kinases (CDKs) regulate a variety of fundamental cellular processes. CDK10 stands out as one of the last orphan CDKs for which no activating cyclin has been identified and no kinase activity revealed. Previous work has shown that CDK10 silencing increases ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2)-driven activation of the MAPK pathway, which confers tamoxifen resistance to breast cancer cells. The precise mechanisms by which CDK10 modulates ETS2 activity, and more generally the functions of CDK10, remain elusive. Here we demonstrate that CDK10 is a cyclin-dependent kinase by identifying cyclin M as an activating cyclin. Cyclin M, an orphan cyclin, is the product of FAM58A, whose mutations cause STAR syndrome, a human developmental anomaly whose features include toe syndactyly, telecanthus, and anogenital and renal malformations. We show that STAR syndrome-associated cyclin M mutants are unable to interact with CDK10. Cyclin M silencing phenocopies CDK10 silencing in increasing c-Raf and in conferring tamoxifen resistance to breast cancer cells. CDK10/cyclin M phosphorylates ETS2 in vitro, and in cells it positively controls ETS2 degradation by the proteasome. ETS2 protein levels are increased in cells derived from a STAR patient, and this increase is attributable to decreased cyclin M levels. Altogether, our results reveal an additional regulatory mechanism for ETS2, which plays key roles in cancer and development. They also shed light on the molecular mechanisms underlying STAR syndrome. Cyclin-dependent kinases (CDKs) play a pivotal role in the control of a number of fundamental cellular processes (1). The human genome contains 21 genes encoding proteins that can be considered as members of the CDK family owing to their sequence similarity with bona fide CDKs, those known to be activated by cyclins (2). Although discovered almost 20 y ago (3, 4), CDK10 remains one of the two CDKs without an identified cyclin partner. This knowledge gap has largely impeded the exploration of its biological functions. CDK10 can act as a positive cell cycle regulator in some cells (5, 6) or as a tumor suppressor in others (7, 8). CDK10 interacts with the ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2) transcription factor and inhibits its transcriptional activity through an unknown mechanism (9). CDK10 knockdown derepresses ETS2, which increases the expression of the c-Raf protein kinase, activates the MAPK pathway, and induces resistance of MCF7 cells to tamoxifen (6). ...

**2.2. Mapping the real-world problem to an ML problem****2.2.1. Type of Machine Learning Problem**

There are nine different classes a genetic mutation can be classified into => Multi class classification problem

### 2.2.2. Performance Metric

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation> (<https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation>)

Metric(s):

- Multi class log-loss
- Confusion matrix

### 2.2.3. Machine Learning Objectives and Constraints

Objective: Predict the probability of each data-point belonging to each of the nine classes.

Constraints:

- Interpretability
- Class probabilities are needed.
- Penalize the errors in class probabilities => Metric is Log-loss.
- No Latency constraints.

## 2.3. Train, CV and Test Datasets

Split the dataset randomly into three parts train, cross validation and test with 64%, 16%, 20% of data respectively

## 3. Exploratory Data Analysis

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
```

D:\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:144: FutureWarning: The sklearn.metrics.classification module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.metrics. Anything that cannot be imported from sklearn.metrics is now part of the private API.

warnings.warn(message, FutureWarning)

D:\Anaconda3\lib\site-packages\sklearn\externals\six.py:31: FutureWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (<https://pypi.org/project/six/>).

"(<https://pypi.org/project/six/>).", FutureWarning)

D:\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:144: FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.neighbors. Anything that cannot be imported from sklearn.neighbors is now part of the private API.

warnings.warn(message, FutureWarning)

D:\Anaconda3\lib\site-packages\sklearn\externals\joblib\\_\_init\_\_.py:15: FutureWarning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please import this functionality directly from joblib, which can be installed with: pip install joblib. If this warning is raised when loading pickled models, you may need to re-serialize those models with scikit-learn 0.21+.

warnings.warn(msg, category=FutureWarning)

## 3.1. Reading Data

### 3.1.1. Reading Gene and Variation Data

```
In [2]: data = pd.read_csv('training/training_variants')
print('Number of data points : ', data.shape[0])
print('Number of features : ', data.shape[1])
print('Features : ', data.columns.values)
data.head()
```

Number of data points : 3321

Number of features : 4

Features : ['ID' 'Gene' 'Variation' 'Class']

Out[2]:

	ID	Gene	Variation	Class
0	0	FAM58A	Truncating Mutations	1
1	1	CBL	W802*	2
2	2	CBL	Q249E	2
3	3	CBL	N454D	3
4	4	CBL	L399V	4

```
In [3]: data.shape
```

```
Out[3]: (3321, 4)
```

training/training\_variants is a comma separated file containing the description of the genetic mutations used for training.

Fields are

- **ID** : the id of the row used to link the mutation to the clinical evidence
- **Gene** : the gene where this genetic mutation is located
- **Variation** : the aminoacid change for this mutations
- **Class** : 1-9 the class this genetic mutation has been classified on

### 3.1.2. Reading Text Data

```
In [4]: # note the separator in this file
data_text = pd.read_csv("training/training_text", sep="\\|\\|", engine="python", names=["ID", "TEXT"], skiprows=1)
print('Number of data points : ', data_text.shape[0])
print('Number of features : ', data_text.shape[1])
print('Features : ', data_text.columns.values)
data_text.head()
```

```
Number of data points : 3321
Number of features : 2
Features : ['ID' 'TEXT']
```

```
Out[4]:
```

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...
3	3	Recent evidence has demonstrated that acquired...
4	4	Oncogenic mutations in the monomeric Casitas B...

### 3.1.3. Preprocessing of text

```
In [5]: # Loading stop words from nltk library
stop_words = set(stopwords.words('english'))

def nlp_preprocessing(total_text, index, column):
    if type(total_text) is not int:
        string = ""
        # replace every special char with space
        total_text = re.sub('[^a-zA-Z0-9\n]', ' ', total_text)
        # replace multiple spaces with single space
        total_text = re.sub('\s+', ' ', total_text)
        # converting all the chars into lower-case.
        total_text = total_text.lower()

        for word in total_text.split():
            # if the word is a not a stop word then retain that word from the data
            if not word in stop_words:
                string += word + " "

        data_text[column][index] = string
```

```
In [6]: #text processing stage.
start_time = time.clock()
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        nlp_preprocessing(row['TEXT'], index, 'TEXT')
    else:
        print("there is no text description for id:",index)
print('Time took for preprocessing the text :',time.clock() - start_time, "seconds")
```

```
there is no text description for id: 1109
there is no text description for id: 1277
there is no text description for id: 1407
there is no text description for id: 1639
there is no text description for id: 2755
Time took for preprocessing the text : 220.5662871749492 seconds
```



```
In [7]: #merging both gene_variations and text data based on ID
result = pd.merge(data, data_text,on='ID', how='left')
result.head()
```

Out[7]:

	ID	Gene	Variation	Class	TEXT
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...
1	1	CBL	W802*	2	abstract background non small cell lung cancer...
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...

```
In [8]: result[result.isnull().any(axis=1)]
```

Out[8]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	NaN
1277	1277	ARID5B	Truncating Mutations	1	NaN
1407	1407	FGFR3	K508M	6	NaN
1639	1639	FLT1	Amplification	6	NaN
2755	2755	BRAF	G596C	7	NaN

```
In [9]: result.loc[result['TEXT'].isnull(),'TEXT'] = result['Gene'] + ' '+result['Variation']
```

```
In [10]: result[result['ID']==1109]
```

Out[10]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	FANCA S1088F

### 3.1.4. Test, Train and Cross Validation Split

#### 3.1.4.1. Splitting data into train, test and cross validation (64:20:16)

```
In [11]: y_true = result['Class'].values
result.Gene      = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')

# split the data into test and train by maintaining same distribution of output variable 'y_true' [stratify=y_true]
X_train, test_df, y_train, y_test = train_test_split(result, y_true, stratify=y_true, test_size=0.2)
# split the train data into train and cross validation by maintaining same distribution of output variable 'y_train' [stratify=y_train]
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, stratify=y_train, test_size=0.2)
```

We split the data into train, test and cross validation data sets, preserving the ratio of class distribution in the original data set

```
In [12]: print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])
```

```
Number of data points in train data: 2124
Number of data points in test data: 665
Number of data points in cross validation data: 532
```

### 3.1.4.2. Distribution of y\_i's in Train, Test and Cross Validation datasets

```

In [13]: # it returns a dict, keys as class labels and values as the number of data points in that class
train_class_distribution = train_df['Class'].value_counts().sortlevel()
test_class_distribution = test_df['Class'].value_counts().sortlevel()
cv_class_distribution = cv_df['Class'].value_counts().sortlevel()

my_colors = 'rgbkymc'
train_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in train data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', train_class_distribution.values[i], '(', np.round((train_class_distribution.values[i]/train_df.shape[0]*100), 3), '%)')

print('-'*80)
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
plt.show()


# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-test_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', test_class_distribution.values[i], '(', np.round((test_class_distribution.values[i]/test_df.shape[0]*100), 3), '%)')

print('-'*80)
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order


```

```
g order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', cv_class_distribution.values[i], '(', np.round((cv_class_distribution.values[i]/cv_df.shape[0]*100), 3), '%)')
```




Number of data points in class 7 : 609 ( 28.672 %)  
Number of data points in class 4 : 439 ( 20.669 %)  
Number of data points in class 1 : 363 ( 17.09 %)  
Number of data points in class 2 : 289 ( 13.606 %)  
Number of data points in class 6 : 176 ( 8.286 %)  
Number of data points in class 5 : 155 ( 7.298 %)  
Number of data points in class 3 : 57 ( 2.684 %)  
Number of data points in class 9 : 24 ( 1.13 %)  
Number of data points in class 8 : 12 ( 0.565 %)

-----  
---



Number of data points in class 7 : 191 ( 28.722 %)  
Number of data points in class 4 : 137 ( 20.602 %)  
Number of data points in class 1 : 114 ( 17.143 %)  
Number of data points in class 2 : 91 ( 13.684 %)  
Number of data points in class 6 : 55 ( 8.271 %)  
Number of data points in class 5 : 48 ( 7.218 %)  
Number of data points in class 3 : 18 ( 2.707 %)  
Number of data points in class 9 : 7 ( 1.053 %)  
Number of data points in class 8 : 4 ( 0.602 %)

-----  
---



Number of data points in class 7 : 153 ( 28.759 %)  
Number of data points in class 4 : 110 ( 20.677 %)  
Number of data points in class 1 : 91 ( 17.105 %)  
Number of data points in class 2 : 72 ( 13.534 %)  
Number of data points in class 6 : 44 ( 8.271 %)  
Number of data points in class 5 : 39 ( 7.331 %)  
Number of data points in class 3 : 14 ( 2.632 %)  
Number of data points in class 9 : 6 ( 1.128 %)  
Number of data points in class 8 : 3 ( 0.564 %)

## 3.2 Prediction using a 'Random' Model

In a 'Random' Model, we generate the NINE class probabilities randomly such that they sum to 1.

```

In [14]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i a
    re predicted class j

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in th
    at column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 correesonds to columns and axis=1 corresponds to
    rows in two dimensional array
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in th
    at row

    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 correesonds to columns and axis=1 corresponds to
    rows in two dimensional array
    # C.sum(axix =0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]

    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, y
    ticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

    print("-"*20, "Precision matrix (Columm Sum=1)", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, y
    ticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

    # representing B in heatmap format
    print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
    plt.figure(figsize=(20,7))

```

```
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, y  
ticklabels=labels)  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
In [15]: # we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y, eps=1e-15))

# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, eps=1e-15))

predicted_y =np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)
```


Log loss on Cross Validation Data using Random Model 2.5062271926893938

Log loss on Test Data using Random Model 2.4869298665993456


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 3.3 Univariate Analysis



```

In [16]: # code for response coding with Laplace smoothing.
# alpha : used for Laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurrences of given feature in
# train data dataframe
# build a vector (1*9) , the first element = (number of times it occurred in cl
# ass1 + 10*alpha / number of time it occurred in total data+90*alpha)
# gv_dict is like a look up table, for every gene it store a (1*9) representat
# ion of it
# for a value of feature in df:
# if it is in train data:
# we add the vector that was stored in 'gv_dict' look up table to 'gv_fea'
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9,1/9, 1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----

# get_gv_fea_dict: Get Gene variation Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
    #      {BRCA1      174
    #       TP53      106
    #       EGFR       86
    #       BRCA2       75
    #       PTEN       69
    #       KIT        61
    #       BRAF       60
    #       ERBB2       47
    #       PDGFRA      46
    #       ...}
    # print(train_df['Variation'].value_counts())
    # output:
    # {
    #   Truncating_Mutations      63
    #   Deletion                   43
    #   Amplification              43
    #   Fusions                    22
    #   Overexpression             3
    #   E17K                       3
    #   Q61L                       3
    #   S222D                      2
    #   P130S                      2
    #   ...
    # }
    value_count = train_df[feature].value_counts()

    # gv_dict : Gene Variation Dict, which contains the probability array for
    # each gene/variation
    gv_dict = dict()

    # denominator will contain the number of time that particular feature occu

```

```

red in whole data
    for i, denominator in value_count.items():
        # vec will contain (p(yi=1/Gi) probability of gene/variation belongs
        to particular class
        # vec is 9 dimensional vector
        vec = []
        for k in range(1,10):
            # print(train_df.loc[(train_df['Class']==1) & (train_df['Gene']=
            = 'BRCA1')])
            #
            ID      Gene      Variation      Class
            # 2470  2470  BRCA1      S1715C      1
            # 2486  2486  BRCA1      S1841R      1
            # 2614  2614  BRCA1      M1R      1
            # 2432  2432  BRCA1      L1657P      1
            # 2567  2567  BRCA1      T1685A      1
            # 2583  2583  BRCA1      E1660G      1
            # 2634  2634  BRCA1      W1718L      1
            # cls_cnt.shape[0] will return the number of rows

            cls_cnt = train_df.loc[(train_df['Class']==k) & (train_df[feature]
            ==i)]

            # cls_cnt.shape[0](numerator) will contain the number of time that
            particular feature occurred in whole data
            vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 90*alpha
            ))

            # we are adding the gene/variation to the dict as key and vec as value
            gv_dict[i]=vec
        return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    # {'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.06818181818181818, 0.13636363636363635, 0.25, 0.19318181818181818, 0.03787878787878788, 0.03787878787878788, 0.03787878787878788],
    # 'TP53': [0.32142857142857145, 0.061224489795918366, 0.061224489795918366, 0.27040816326530615, 0.061224489795918366, 0.066326530612244902, 0.051020408163265307, 0.051020408163265307, 0.056122448979591837],
    # 'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.068181818181818177, 0.068181818181818177, 0.0625, 0.34659090909090912, 0.0625, 0.056818181818181816],
    # 'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.060606060606060608, 0.078787878787878782, 0.139393939393939394, 0.34545454545454546, 0.060606060606060608, 0.060606060606060608, 0.060606060606060608],
    # 'PTEN': [0.069182389937106917, 0.062893081761006289, 0.069182389937106917, 0.46540880503144655, 0.075471698113207544, 0.062893081761006289, 0.069182389937106917, 0.062893081761006289, 0.062893081761006289],
    # 'KIT': [0.066225165562913912, 0.25165562913907286, 0.072847682119205295, 0.072847682119205295, 0.066225165562913912, 0.066225165562913912, 0.27152317880794702, 0.066225165562913912, 0.066225165562913912],
    # 'BRAF': [0.066666666666666666, 0.17999999999999999, 0.07333333333333333, 0.073333333333333334, 0.093333333333333338, 0.080000000000000002, 0.29999999999999999, 0.066666666666666666, 0.066666666666666666],
    # ...
    # }

```

```

gv_dict = get_gv_fea_dict(alpha, feature, df)
# value_count is similar in get_gv_fea_dict
value_count = train_df[feature].value_counts()

# gv_fea: Gene_variation feature, it will contain the feature for each fea
ture value in the data
gv_fea = []
# for every feature values in the given data frame we will check if it is
there in the train data then we will add the feature to gv_fea
# if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
for index, row in df.iterrows():
    if row[feature] in dict(value_count).keys():
        gv_fea.append(gv_dict[row[feature]])
    else:
        gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
#
return gv_fea

```

when we caculate the probability of a feature belongs to any particular class, we apply laplace smoothing

- $(\text{numerator} + 10 \cdot \alpha) / (\text{denominator} + 90 \cdot \alpha)$

### 3.2.1 Univariate Analysis on Gene Feature

**Q1.** Gene, What type of feature it is ?

**Ans.** Gene is a categorical variable

**Q2.** How many categories are there and How they are distributed?

```

In [17]: unique_genes = train_df['Gene'].value_counts()
print('Number of Unique Genes :', unique_genes.shape[0])
# the top 10 genes that occured most
print(unique_genes.head(10))

```

```


Number of Unique Genes : 225
BRCA1      173
TP53       108
EGFR       100
PTEN        85
BRCA2       79
KIT         58
BRAF        47
PDGFRA      44
ERBB2       44
ALK         43
Name: Gene, dtype: int64

```


```
In [18]: print("Ans: There are", unique_genes.shape[0] ,"different categories of genes  
in the train data, and they are distributed as follows",)
```

Ans: There are 225 different categories of genes in the train data, and they are distributed as follows

```
In [19]: s = sum(unique_genes.values);  
h = unique_genes.values/s;  
plt.plot(h, label="Histogram of Genes")  
plt.xlabel('Index of a Gene')  
plt.ylabel('Number of Occurances')  
plt.legend()  
plt.grid()  
plt.show()
```



```
In [20]: c = np.cumsum(h)  
plt.plot(c, label='Cumulative distribution of Genes')  
plt.grid()  
plt.legend()  
plt.show()
```



### Q3. How to featurize this Gene feature ?

**Ans.**there are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will choose the appropriate featurization based on the ML model we use. For this problem of multi-class classification with categorical features, one-hot encoding is better for Logistic regression while response coding is better for Random Forests.

```
In [21]: #response-coding of the Gene feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", train_df))
# test gene feature
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", test_df))
# cross validation gene feature
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))
```

```
In [22]: print("train_gene_feature_responseCoding is converted feature using response coding method. The shape of gene feature:", train_gene_feature_responseCoding.shape)
print("test_gene_feature_responseCoding is converted feature using response coding method. The shape of test feature:", test_gene_feature_responseCoding.shape)
print("cv_gene_feature_responseCoding is converted feature using response coding method. The shape of cv feature:", cv_gene_feature_responseCoding.shape)
```

```
train_gene_feature_responseCoding is converted feature using response coding method. The shape of gene feature: (2124, 9)
test_gene_feature_responseCoding is converted feature using response coding method. The shape of test feature: (665, 9)
cv_gene_feature_responseCoding is converted feature using response coding method. The shape of cv feature: (532, 9)
```

```
In [23]: # one-hot encoding of Gene feature.
gene_vectorizer = TfidfVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```



```
In [24]: train_df['Gene'].head()
```

```
Out[24]: 2998      KIT  
2848      BRCA2  
1799      AR  
383      TP53  
2722      BRAF  
Name: Gene, dtype: object
```

```
In [25]: gene_vectorizer.get_feature_names()
```

```
Out[25]: ['abl1',  
          'acvr1',  
          'ago2',  
          'akt1',  
          'akt2',  
          'akt3',  
          'alk',  
          'apc',  
          'ar',  
          'araf',  
          'arid1a',  
          'arid1b',  
          'arid2',  
          'arid5b',  
          'asx11',  
          'asx12',  
          'atm',  
          'atr',  
          'atrx',  
          'aurka',  
          'aurkb',  
          'axin1',  
          'b2m',  
          'bap1',  
          'bard1',  
          'bcl10',  
          'bcl2l11',  
          'bcor',  
          'braf',  
          'brca1',  
          'brca2',  
          'brd4',  
          'brip1',  
          'btk',  
          'card11',  
          'carm1',  
          'casp8',  
          'cbl',  
          'ccnd1',  
          'ccnd2',  
          'ccnd3',  
          'cdh1',  
          'cdk12',  
          'cdk4',  
          'cdk6',  
          'cdk8',  
          'cdkn1a',  
          'cdkn1b',  
          'cdkn2a',  
          'cdkn2b',  
          'cdkn2c',  
          'chek2',  
          'cic',  
          'crebbp',  
          'ctcf',  
          'ctnnb1',  
          'ddr2',
```

'dicer1',  
'dnmt3a',  
'dnmt3b',  
'egfr',  
'eif1ax',  
'elf3',  
'ep300',  
'epas1',  
'epcam',  
'erbb2',  
'erbb3',  
'erbb4',  
'ercc2',  
'ercc4',  
'erg',  
'esr1',  
'etv1',  
'etv6',  
'ewsr1',  
'ezh2',  
'fanca',  
'fancc',  
'fat1',  
'fbxw7',  
'fgf19',  
'fgf3',  
'fgfr1',  
'fgfr2',  
'fgfr3',  
'fgfr4',  
'flt3',  
'foxa1',  
'foxp1',  
'gata3',  
'gli1',  
'gna11',  
'gnas',  
'h3f3a',  
'hist1h1c',  
'hla',  
'hnf1a',  
'hras',  
'idh1',  
'idh2',  
'igf1r',  
'il7r',  
'inpp4b',  
'jak1',  
'jak2',  
'kdm5a',  
'kdm5c',  
'kdm6a',  
'kdr',  
'keap1',  
'kit',  
'kmt2a',  
'kmt2b',

'kmt2c',  
'kmt2d',  
'knstrn',  
'kras',  
'lats2',  
'map2k1',  
'map2k2',  
'map2k4',  
'map3k1',  
'mapk1',  
'med12',  
'mef2b',  
'met',  
'mga',  
'mlh1',  
'mpl',  
'msh2',  
'msh6',  
'mtor',  
'myc',  
'mycn',  
'myd88',  
'ncor1',  
'nf1',  
'nf2',  
'nfe2l2',  
'nfkb1a',  
'nkx2',  
'notch1',  
'notch2',  
'npm1',  
'nras',  
'nsd1',  
'ntrk1',  
'ntrk2',  
'nup93',  
'pak1',  
'pax8',  
'pdgfra',  
'pdgfrb',  
'pik3ca',  
'pik3cb',  
'pik3cd',  
'pik3r1',  
'pik3r2',  
'pim1',  
'pms1',  
'pms2',  
'pole',  
'ppm1d',  
'ppp2r1a',  
'ppp6c',  
'prdm1',  
'ptch1',  
'pten',  
'ptpn11',  
'ptprd',

```
'ptprt',  
'rab35',  
'rac1',  
'rad21',  
'rad50',  
'rad51b',  
'rad51c',  
'rad51d',  
'rad54l',  
'raf1',  
'rasa1',  
'rb1',  
'rbm10',  
'ret',  
'rheb',  
'rhoa',  
'rit1',  
'ros1',  
'runx1',  
'rxra',  
'sdhc',  
'sf3b1',  
'shq1',  
'smad2',  
'smad3',  
'smad4',  
'smarca4',  
'smarcb1',  
'smo',  
'sos1',  
'sox9',  
'spop',  
'src',  
'srsf2',  
'stag2',  
'stat3',  
'stk11',  
'tert',  
'tet2',  
'tgfbr1',  
'tgfbr2',  
'tmprss2',  
'tp53',  
'tp53bp1',  
'tsc1',  
'tsc2',  
'u2af1',  
'vhl',  
'whsc1',  
'whsc1l1',  
'xpo1',  
'xrcc2',  
'yap1']
```

```
In [26]: print("train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature:", train_gene_feature_onehotCoding.shape)
```

train\_gene\_feature\_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature: (2124, 224)

#### Q4. How good is this gene feature in predicting $y_i$ ?

There are many ways to estimate how good a feature is, in predicting  $y_i$ . One of the good methods is to build a proper ML model using just this feature. In this case, we will build a logistic regression model using only Gene feature (one hot encoded) to predict  $y_i$ .

```

In [27]: alpha = [10 ** x for x in range(-5, 1)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_gene_feature_onehotCoding, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_gene_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_gene_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)


predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)

```



```
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss i
s:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

For values of alpha = 1e-05 The log loss is: 1.2540111267947311  
 For values of alpha = 0.0001 The log loss is: 1.2120950763061806  
 For values of alpha = 0.001 The log loss is: 1.239000589887041  
 For values of alpha = 0.01 The log loss is: 1.3186440894440317  
 For values of alpha = 0.1 The log loss is: 1.408721964712307  
 For values of alpha = 1 The log loss is: 1.4553525134680059



For values of best alpha = 0.0001 The train log loss is: 0.9858999925424323  
 For values of best alpha = 0.0001 The cross validation log loss is: 1.2120950763061806  
 For values of best alpha = 0.0001 The test log loss is: 1.1945034241197496

**Q5.** Is the Gene feature stable across all the data sets (Test, Train, Cross validation)?

**Ans.** Yes, it is. Otherwise, the CV and Test errors would be significantly more than train error.

```
In [28]: print("Q6. How many data points in Test and CV datasets are covered by the ",
unique_genes.shape[0], " genes in train dataset?")

test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]

print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0], ":",(te
st_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],":",
(cv_coverage/cv_df.shape[0])*100)
```

Q6. How many data points in Test and CV datasets are covered by the 225 genes in train dataset?

Ans

1. In test data 638 out of 665 : 95.93984962406014
2. In cross validation data 507 out of 532 : 95.30075187969925

### 3.2.2 Univariate Analysis on Variation Feature

**Q7.** Variation, What type of feature is it ?

**Ans.** Variation is a categorical variable

**Q8.** How many categories are there?


```
In [29]: unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
# the top 10 variations that occurred most
print(unique_variations.head(10))
```

```
Number of Unique Variations : 1939
Truncating_Mutations      52
Deletion                   50
Amplification              41
Fusions                    15
Overexpression             4
Q61L                       3
Q61R                       3
Y64A                       2
R170W                      2
E330K                      2
Name: Variation, dtype: int64
```

```
In [30]: print("Ans: There are", unique_variations.shape[0] , "different categories of v
          ariations in the train data, and they are distributed as follows",)
```


```
Ans: There are 1939 different categories of variations in the train data, and
they are distributed as follows
```

```
In [31]: s = sum(unique_variations.values);
h = unique_variations.values/s;
plt.plot(h, label="Histogram of Variations")
plt.xlabel('Index of a Variation')
plt.ylabel('Number of Occurances')
plt.legend()
plt.grid()
plt.show()
```



```
In [32]: c = np.cumsum(h)
print(c)
plt.plot(c, label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()
```

```
[0.02448211 0.0480226 0.0673258 ... 0.99905838 0.99952919 1. ]
```



## Q9. How to featurize this Variation feature ?

**Ans.** There are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will be using both these methods to featurize the Variation Feature

```
In [33]: # alpha is used for Laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", train_df))
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", test_df))
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_df))
```

```
In [34]: print("train_variation_feature_responseCoding is a converted feature using the
response coding method. The shape of Variation feature:", train_variation_feature_responseCoding.shape)
```

train\_variation\_feature\_responseCoding is a converted feature using the response coding method. The shape of Variation feature: (2124, 9)

```
In [35]: # one-hot encoding of variation feature.
variation_vectorizer = TfidfVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])
```

```
In [36]: print("train_variation_feature_onehotEncoded is converted feature using the one-hot
encoding method. The shape of Variation feature:", train_variation_feature_onehotCoding.shape)
```

train\_variation\_feature\_onehotEncoded is converted feature using the one-hot encoding method. The shape of Variation feature: (2124, 1973)

## Q10. How good is this Variation feature in predicting $y_i$ ?

Let's build a model just like the earlier!

```

In [37]: alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_variation_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_variation_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)

    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_variation_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is")

```

```
s:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss i
s:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

For values of alpha = 1e-05 The log loss is: 1.6955477942667168  
 For values of alpha = 0.0001 The log loss is: 1.6903773117608554  
 For values of alpha = 0.001 The log loss is: 1.6965718919110084  
 For values of alpha = 0.01 The log loss is: 1.7126152285050742  
 For values of alpha = 0.1 The log loss is: 1.7121025941256351  
 For values of alpha = 1 The log loss is: 1.7135711277220442



For values of best alpha = 0.0001 The train log loss is: 0.6900581252091896  
 For values of best alpha = 0.0001 The cross validation log loss is: 1.6903773117608554  
 For values of best alpha = 0.0001 The test log loss is: 1.7027076929439666

**Q11.** Is the Variation feature stable across all the data sets (Test, Train, Cross validation)?

**Ans.** Not sure! But lets be very sure using the below analysis.

```
In [38]: print("Q12. How many data points are covered by total ", unique_variations.shape[0], " genes in test and cross validation data sets?")
test_coverage=test_df[test_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0], ":",(test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],":", (cv_coverage/cv_df.shape[0])*100)
```

Q12. How many data points are covered by total 1939 genes in test and cross validation data sets?

Ans

1. In test data 76 out of 665 : 11.428571428571429

2. In cross validation data 59 out of 532 : 11.090225563909774

### 3.2.3 Univariate Analysis on Text Feature

1. How many unique words are present in train data?
2. How are word frequencies distributed?
3. How to featurize text field?
4. Is the text feature useful in predicting  $y_i$ ?
5. Is the text feature stable across train, test and CV datasets?

```
In [39]: # cls_text is a data frame
# for every row in data frame consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word

def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary
```

```
In [40]: import math
#https://stackoverflow.com/a/1602964
def get_text_responsecoding(df):
    text_feature_responseCoding = np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index = 0
        for index, row in df.iterrows():
            sum_prob = 0
            for word in row['TEXT'].split():
                sum_prob += math.log(((dict_list[i].get(word,0)+10 )/(total_dict.get(word,0)+90)))
            text_feature_responseCoding[row_index][i] = math.exp(sum_prob/len(row['TEXT'].split()))
            row_index += 1
    return text_feature_responseCoding
```

```
In [41]: # building a CountVectorizer with all the words that occurred minimum 3 times in train data
text_vectorizer = TfidfVectorizer(min_df=5, max_features=3000)
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and return s (1*number of features) vector
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its number of times it occurred
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_counts))

print("Total number of unique words in train data :", len(train_text_features))
```

Total number of unique words in train data : 3000



```

In [42]: dict_list = []
# dict_list =[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list

# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dictionary_paddle(train_df)

confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)

```

```

In [43]: #response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)

```

```

In [44]: # https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding = (train_text_feature_responseCoding.T/train_text_feature_responseCoding.sum(axis=1)).T
test_text_feature_responseCoding = (test_text_feature_responseCoding.T/test_text_feature_responseCoding.sum(axis=1)).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T/cv_text_feature_responseCoding.sum(axis=1)).T

```

```

In [45]: # don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)

```

```
In [46]: #https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1] ,
reverse=True))
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))
```

```
In [47]: # Number of words for a given frequency.  
print(Counter(sorted_text_occur))
```

Counter({193.87177083533788: 1, 127.28309351290949: 1, 123.9516369205886: 1, 95.27634367108831: 1, 89.49135585896761: 1, 88.45930440253197: 1, 87.76525239 990544: 1, 82.71406689968998: 1, 82.30785911801122: 1, 80.93104160221537: 1, 78.18050812460085: 1, 76.53874053412723: 1, 73.26940924075917: 1, 73.11228364 820379: 1, 68.55171974419262: 1, 62.968953418572504: 1, 59.55562691755885: 1, 59.510612038449665: 1, 59.1312029163687: 1, 58.710685016121815: 1, 54.9615870 39993695: 1, 53.50305936525071: 1, 52.75427793736526: 1, 50.009559714074165: 1, 49.319185747759235: 1, 48.574687892104016: 1, 48.431909215092396: 1, 47.85 8275167061656: 1, 47.29853833934683: 1, 47.09753485344249: 1, 46.685476140911 18: 1, 46.633634991667: 1, 46.39696865552379: 1, 46.12490703291477: 1, 44.349 5023026693: 1, 43.75740487272057: 1, 42.39872058071737: 1, 42.04550958731254: 1, 41.852064292591606: 1, 41.69793644078394: 1, 41.236905830027645: 1, 38.768 32817970881: 1, 38.36146652266362: 1, 38.28310233781683: 1, 38.18379421696625 5: 1, 37.19653306267872: 1, 36.73446706081922: 1, 35.35206271696007: 1, 35.30 4245996714904: 1, 34.827951309150414: 1, 34.780114553426515: 1, 34.0698445848 9117: 1, 33.92291938153465: 1, 33.9041744517319: 1, 33.89854120904206: 1, 32. 942589686711024: 1, 32.737986458638595: 1, 32.7064445822106: 1, 32.6018996433 4747: 1, 32.34950914849501: 1, 32.15076175713733: 1, 31.84663263110182: 1, 3 1.816163172959843: 1, 31.537507132038073: 1, 31.465391384465065: 1, 31.404846 060824234: 1, 31.179353318734737: 1, 31.04579669688957: 1, 30.93285688129221: 1, 30.88685889973828: 1, 30.384306328324268: 1, 30.025024180905522: 1, 29.770 49631673016: 1, 29.05573376031132: 1, 28.64883117005774: 1, 28.06417342315510 4: 1, 28.010170657849454: 1, 27.970425044562244: 1, 27.820646515406946: 1, 2 7.803149904907784: 1, 27.182720099938887: 1, 27.149771188392343: 1, 26.727076 375650906: 1, 26.517874815548407: 1, 26.38186325191481: 1, 26.14252933000759: 1, 26.115156805921895: 1, 25.79741421789772: 1, 25.639982812347817: 1, 25.611 866816098853: 1, 25.584730069851364: 1, 25.536882700841442: 1, 25.32969356052 2813: 1, 25.29846431435189: 1, 25.29668307254763: 1, 25.236628476995605: 1, 2 4.55330129720298: 1, 24.494726724436862: 1, 24.341945191668792: 1, 24.1582531 1610082: 1, 24.14330860512465: 1, 24.12679440864085: 1, 24.086753080561554: 1, 23.82946234270497: 1, 23.821917601049005: 1, 23.712575922538772: 1, 23.545 73826737324: 1, 23.47870478753906: 1, 23.387292663136407: 1, 23.3775328814707 07: 1, 23.340693666195193: 1, 23.268320701511172: 1, 23.18263950738593: 1, 2 3.155641711161437: 1, 22.950608103298716: 1, 22.857600219823762: 1, 22.813696 116894942: 1, 22.77032918483672: 1, 22.675767248823774: 1, 22.64985415296104 4: 1, 22.645705311908795: 1, 22.586473868991188: 1, 22.36137565698606: 1, 22. 287709101263392: 1, 22.09056389506355: 1, 22.00807404739081: 1, 21.6787732558 21426: 1, 21.640104608349674: 1, 21.584191346543385: 1, 21.577724201250284: 1, 21.52389504676622: 1, 21.470735472247632: 1, 21.444506386428763: 1, 21.418 982053153037: 1, 21.395844854374257: 1, 21.359828044008662: 1, 21.27975089268 496: 1, 21.223390604842685: 1, 21.19860023804342: 1, 20.851213292892687: 1, 2 0.745877206617475: 1, 20.52753789104859: 1, 20.523727758778954: 1, 20.4891284 8544538: 1, 20.33910766828922: 1, 20.167470548830185: 1, 20.154769057058118: 1, 20.06852057701836: 1, 20.03566976446402: 1, 19.862797219185435: 1, 19.8451 9935472838: 1, 19.76158558931519: 1, 19.742694582458224: 1, 19.64712471458678 4: 1, 19.46151857851957: 1, 19.24653259411346: 1, 19.245469968784846: 1, 19.1 54599758222854: 1, 19.14895836489851: 1, 19.112079013412984: 1, 19.0459202834 70775: 1, 18.960873782331557: 1, 18.93364273704337: 1, 18.721558286760274: 1, 18.62943376782954: 1, 18.595130187655258: 1, 18.44946480649205: 1, 18.4111557 86653406: 1, 18.359536652868297: 1, 18.35848275111671: 1, 18.248329827819315: 1, 18.09490245683486: 1, 18.084199487822094: 1, 18.044819361193397: 1, 17.990 12332399031: 1, 17.949204085379215: 1, 17.934923217489583: 1, 17.884234630584 78: 1, 17.872702364320283: 1, 17.74758945302231: 1, 17.703948332579003: 1, 1 7.637613373092826: 1, 17.585836455695883: 1, 17.56794413494214: 1, 17.4886831 3947405: 1, 17.45907203010106: 1, 17.38980471286387: 1, 17.335336445756194: 1, 17.317338515621095: 1, 17.287297644762898: 1, 17.274407813691397: 1, 17.22 8776800765267: 1, 17.189025493079217: 1, 17.080053051256744: 1, 17.0269961198

17124: 1, 16.94736850576957: 1, 16.937238260697242: 1, 16.92707837740094: 1, 16.873267800376592: 1, 16.80757203255165: 1, 16.788755520065184: 1, 16.787797607882883: 1, 16.695541756564026: 1, 16.65966348913304: 1, 16.60687298507159: 1, 16.552940487685934: 1, 16.534108090874952: 1, 16.52738883406207: 1, 16.51414114554796: 1, 16.50707629020627: 1, 16.48939477083588: 1, 16.482009196980826: 1, 16.466728204503095: 1, 16.432979680189405: 1, 16.37123728324817: 1, 16.35748968600242: 1, 16.356882471162866: 1, 16.32390347519769: 1, 16.31835756951897: 1, 16.30827208639079: 1, 16.2426531365701: 1, 16.18238930489689: 1, 16.148121649469395: 1, 16.079738469387085: 1, 16.022740925202854: 1, 16.015873184019533: 1, 15.966435013976854: 1, 15.932389374362169: 1, 15.886452345572064: 1, 15.876943825690718: 1, 15.818012521863153: 1, 15.808666825687249: 1, 15.788610711518652: 1, 15.756239108108751: 1, 15.753064291926602: 1, 15.736614695904471: 1, 15.706950115728686: 1, 15.671470968163481: 1, 15.637836529658067: 1, 15.633863290073712: 1, 15.504612395888499: 1, 15.490847742885952: 1, 15.366983860938388: 1, 15.364041936815338: 1, 15.195036321084595: 1, 15.135268646984635: 1, 15.111805353268913: 1, 15.086259178431757: 1, 15.05986042962984: 1, 15.057219643825707: 1, 14.981801400176925: 1, 14.968607333865515: 1, 14.955510287433821: 1, 14.88673746281047: 1, 14.849640340755055: 1, 14.818067301214457: 1, 14.814590370005465: 1, 14.813778078036: 1, 14.79891583199587: 1, 14.74178672116593: 1, 14.699812626959647: 1, 14.681173343597512: 1, 14.679841219396309: 1, 14.676736822258105: 1, 14.67551602740085: 1, 14.632358952784674: 1, 14.624541234541802: 1, 14.601375387238047: 1, 14.559617619572581: 1, 14.538656756363382: 1, 14.514724865542027: 1, 14.495833663987867: 1, 14.474689277760081: 1, 14.433408673537075: 1, 14.414712689965071: 1, 14.28274450565617: 1, 14.281013533777983: 1, 14.270527603027809: 1, 14.267767736684812: 1, 14.196975935587318: 1, 14.1446419676323: 1, 14.107024521764437: 1, 14.084751984451666: 1, 14.02654160218381: 1, 13.997834976894534: 1, 13.989710388698853: 1, 13.963819564692901: 1, 13.919518160940248: 1, 13.900835983967696: 1, 13.87159424444936: 1, 13.867320797013845: 1, 13.847494447919305: 1, 13.836576683874544: 1, 13.83563676402577: 1, 13.790796118444796: 1, 13.778248957427495: 1, 13.757861400653118: 1, 13.731206172860306: 1, 13.703502851275656: 1, 13.667795095196828: 1, 13.599589690769248: 1, 13.594575755486794: 1, 13.548688954171345: 1, 13.482762509439539: 1, 13.440569987481874: 1, 13.38641793126872: 1, 13.34258210857564: 1, 13.336525061689692: 1, 13.33123501097349: 1, 13.321021295857292: 1, 13.302693413029436: 1, 13.290895371379934: 1, 13.269384982443809: 1, 13.266528911387299: 1, 13.241690585971504: 1, 13.234784306573644: 1, 13.230600475187822: 1, 13.210330352129608: 1, 13.120336500904813: 1, 13.119645923457323: 1, 13.058778374336736: 1, 13.01087343266874: 1, 12.986081187714483: 1, 12.970309528738015: 1, 12.94979826894049: 1, 12.946334178506508: 1, 12.936812569181653: 1, 12.926583650548002: 1, 12.801458649053924: 1, 12.789286822856653: 1, 12.765275060577721: 1, 12.700110213923047: 1, 12.691891093991586: 1, 12.670334156557304: 1, 12.656930899620283: 1, 12.65454213276898: 1, 12.64067809078919: 1, 12.624107292431814: 1, 12.623569532534193: 1, 12.62355092909542: 1, 12.610246544326051: 1, 12.588405928729927: 1, 12.588283338654781: 1, 12.575372243755382: 1, 12.562889003193655: 1, 12.554073495095567: 1, 12.546541999335133: 1, 12.53586937336321: 1, 12.535129086488594: 1, 12.513274014843175: 1, 12.495926366646575: 1, 12.487044995277369: 1, 12.466053790877998: 1, 12.449490581490972: 1, 12.447726121498402: 1, 12.426755036746556: 1, 12.42291916401745: 1, 12.420311051234448: 1, 12.377034939798671: 1, 12.370828090927223: 1, 12.356617960977506: 1, 12.339517760604481: 1, 12.27494943293379: 1, 12.261163024466233: 1, 12.25476157672538: 1, 12.171768615519705: 1, 12.117801372725683: 1, 12.102088381772747: 1, 12.08056384679612: 1, 12.029243415881265: 1, 12.022783636139009: 1, 12.00654938462864: 1, 11.98564732502804: 1, 11.948078805000627: 1, 11.933705889217519: 1, 11.928296723947097: 1, 11.92624242779842: 1, 11.915946632793808: 1, 11.897180859955334: 1, 11.869265637892836: 1, 11.860519177022093: 1, 11.859561703265424: 1, 11.831308034135919: 1, 11.810038321244033: 1, 11.805611633877767: 1, 11.79884399456977: 1, 11.790977444222223: 1, 11.7576270

4134127: 1, 11.754016871448947: 1, 11.735547316489319: 1, 11.735077016429472: 1, 11.723189394404564: 1, 11.675715424478197: 1, 11.66859802184892: 1, 11.646642557405771: 1, 11.641236555971643: 1, 11.595304163210722: 1, 11.586629693146927: 1, 11.572960605158405: 1, 11.540349087800944: 1, 11.53118849975283: 1, 11.521949347056994: 1, 11.505186603019377: 1, 11.494956703870832: 1, 11.44198830512959: 1, 11.438483693873028: 1, 11.432973388806673: 1, 11.417346884331266: 1, 11.393914347187769: 1, 11.369176165454775: 1, 11.33499683421711: 1, 11.326361259257189: 1, 11.319706476217428: 1, 11.316454184238045: 1, 11.271109886472384: 1, 11.267952043215322: 1, 11.262800379816474: 1, 11.245905539953355: 1, 11.240573424634356: 1, 11.207957405076423: 1, 11.190394513111626: 1, 11.167412210801757: 1, 11.156050346372473: 1, 11.151640515796164: 1, 11.149543641297463: 1, 11.145455750614522: 1, 11.119502412521424: 1, 11.116547945967413: 1, 11.106911752047852: 1, 11.090971580383645: 1, 11.076486282231782: 1, 11.061304284527152: 1, 11.057849570436023: 1, 11.043292235219473: 1, 11.02025595528952: 1, 11.01860704754094: 1, 10.999120840808736: 1, 10.987889596304957: 1, 10.958303770159073: 1, 10.952874416730497: 1, 10.914567431384604: 1, 10.905909991022366: 1, 10.896341608578439: 1, 10.8945934931141: 1, 10.892973110102567: 1, 10.881822260316842: 1, 10.870525389936018: 1, 10.85936594052028: 1, 10.821053850601357: 1, 10.818908333337935: 1, 10.809625350019266: 1, 10.759736751333937: 1, 10.734761028797177: 1, 10.731103387538472: 1, 10.714785044126108: 1, 10.68348802933575: 1, 10.665441499681657: 1, 10.661672298163642: 1, 10.639944025515096: 1, 10.636324378533148: 1, 10.623802604833095: 1, 10.621337228174392: 1, 10.586469343851398: 1, 10.58229441093738: 1, 10.567278968171278: 1, 10.560816531890113: 1, 10.530583110636806: 1, 10.471882122633586: 1, 10.46235628372567: 1, 10.459826244199402: 1, 10.446393346175403: 1, 10.435896614950778: 1, 10.423345042283165: 1, 10.422698826935408: 1, 10.42121626378798: 1, 10.41814743214225: 1, 10.368883307543694: 1, 10.368175039343624: 1, 10.346714753467202: 1, 10.331128244244729: 1, 10.329075385423813: 1, 10.32130378095868: 1, 10.309479215773884: 1, 10.290523067517707: 1, 10.283045284895316: 1, 10.2828971124398: 1, 10.257600836367361: 1, 10.248416371898685: 1, 10.23642738846748: 1, 10.226348908871353: 1, 10.186253314245414: 1, 10.102953708025419: 1, 10.09684000117035: 1, 10.092540605202526: 1, 10.091093325097786: 1, 10.080642747374194: 1, 10.077602284452578: 1, 10.069412761155105: 1, 10.060056088776488: 1, 10.055326130917162: 1, 10.024440861879574: 1, 9.997334858812804: 1, 9.990668467604703: 1, 9.981321137239831: 1, 9.95552819815024: 1, 9.935120376210302: 1, 9.911611034097799: 1, 9.893531407682614: 1, 9.891936001124042: 1, 9.88936108619283: 1, 9.885301376443284: 1, 9.882372727868347: 1, 9.866808429283587: 1, 9.859160429340353: 1, 9.837106209280229: 1, 9.825959420543791: 1, 9.8112448088251: 1, 9.808739008734754: 1, 9.791393898739832: 1, 9.789755911269435: 1, 9.78357253021833: 1, 9.781346474448164: 1, 9.764952805505455: 1, 9.760329807925507: 1, 9.749388520617195: 1, 9.739157942229017: 1, 9.728258610784243: 1, 9.724285438268778: 1, 9.698281630043647: 1, 9.676639661095681: 1, 9.659435616326526: 1, 9.65111566998718: 1, 9.628914668167855: 1, 9.62492077816739: 1, 9.608204622196718: 1, 9.606739990896445: 1, 9.59816570043222: 1, 9.591943325382342: 1, 9.587397265553093: 1, 9.577754950714763: 1, 9.525041199524047: 1, 9.492317166280438: 1, 9.487316487586234: 1, 9.469477587565166: 1, 9.45930136948915: 1, 9.445113479385144: 1, 9.427217202621506: 1, 9.416762999849544: 1, 9.411160665122882: 1, 9.409122141936972: 1, 9.376349999449756: 1, 9.3742191061312: 1, 9.371858026154063: 1, 9.361904995314088: 1, 9.315378950542085: 1, 9.28143027474465: 1, 9.268961027050198: 1, 9.259062366777679: 1, 9.251023038174425: 1, 9.217185599335274: 1, 9.20211545700276: 1, 9.19990619884491: 1, 9.187043445571469: 1, 9.185866448438118: 1, 9.184519227201287: 1, 9.184093434679456: 1, 9.178086407084518: 1, 9.175428826678544: 1, 9.156026353119788: 1, 9.148374535424894: 1, 9.143377390651194: 1, 9.100950231551693: 1, 9.098116449064841: 1, 9.044937083466209: 1, 9.040058643173396: 1, 9.02049783123291: 1, 9.003801261387624: 1, 8.979866547012845: 1, 8.978773525261989: 1, 8.978487724362854: 1, 8.953866282457849: 1, 8.947329022097959: 1, 8.942648688849626: 1, 8.932

77045188741: 1, 8.932665504204666: 1, 8.926632130574177: 1, 8.92568163486072  
1: 1, 8.924023272885329: 1, 8.923405953754594: 1, 8.915137270185033: 1, 8.907  
550727052037: 1, 8.902223625742442: 1, 8.896613872175251: 1, 8.88700961264129  
7: 1, 8.878691966388047: 1, 8.871626292363993: 1, 8.858954799391903: 1, 8.850  
377365589361: 1, 8.848617873776874: 1, 8.820384114231032: 1, 8.81998005070059  
7: 1, 8.809504382352126: 1, 8.807751191632764: 1, 8.784875344464497: 1, 8.757  
974713923772: 1, 8.750481087258304: 1, 8.706195295211645: 1, 8.6870661340181  
4: 1, 8.683674529945973: 1, 8.663113496342492: 1, 8.645339556267366: 1, 8.645  
037333408299: 1, 8.642361272631426: 1, 8.63836389644485: 1, 8.63682229107017  
6: 1, 8.633291329525521: 1, 8.629131975806622: 1, 8.625284497535212: 1, 8.623  
299797175422: 1, 8.622454269864612: 1, 8.539979973995539: 1, 8.53761132018872  
7: 1, 8.526383357673287: 1, 8.499088471221468: 1, 8.497297171883503: 1, 8.490  
587397631803: 1, 8.48005734668533: 1, 8.475042182622056: 1, 8.46057898753323  
7: 1, 8.458639895810867: 1, 8.434965292924764: 1, 8.427282785581891: 1, 8.425  
42285426942: 1, 8.422601676207332: 1, 8.415866259214049: 1, 8.41439689942029  
6: 1, 8.404269383592961: 1, 8.39716354296222: 1, 8.394195124279003: 1, 8.3784  
15984090099: 1, 8.368582925327331: 1, 8.36586746016592: 1, 8.36414558065524:  
1, 8.350669684140657: 1, 8.3469949463602: 1, 8.329227002740948: 1, 8.32333461  
1707082: 1, 8.318904020394863: 1, 8.300499176273261: 1, 8.290087299134651: 1,  
8.287641160770574: 1, 8.286322190307471: 1, 8.263021818951207: 1, 8.259795117  
795553: 1, 8.25312623233653: 1, 8.236973775947327: 1, 8.23673087441342: 1, 8.  
199496202266419: 1, 8.197832515172157: 1, 8.19651904097332: 1, 8.192841185504  
019: 1, 8.189442061573079: 1, 8.178796231610784: 1, 8.176141118068523: 1, 8.1  
59208148532793: 1, 8.158089063196696: 1, 8.153971765921566: 1, 8.145521302184  
696: 1, 8.126870171666573: 1, 8.10922630798033: 1, 8.104947599442662: 1, 8.10  
0345653178847: 1, 8.086735768897473: 1, 8.078165937468501: 1, 8.0758008578393  
92: 1, 8.072893355831859: 1, 8.054767234339854: 1, 8.052612531077932: 1, 8.04  
4135541723708: 1, 8.039082020742235: 1, 8.032799949088858: 1, 8.0234890468528  
23: 1, 8.013688903642468: 1, 8.012721842264762: 1, 8.009694980804325: 1, 8.00  
5382963592792: 1, 7.996739400138342: 1, 7.990426092212076: 1, 7.9901365288587  
69: 1, 7.980166408187319: 1, 7.9754681743928275: 1, 7.974646035199601: 1, 7.9  
703083851853345: 1, 7.965779840409521: 1, 7.959536920143209: 1, 7.94781266533  
9138: 1, 7.93668325716446: 1, 7.9341318952936435: 1, 7.917228985563086: 1, 7.  
9165684536642065: 1, 7.908504131300155: 1, 7.897371299452761: 1, 7.8932922809  
26125: 1, 7.880770532153867: 1, 7.879592339105393: 1, 7.8573833036054515: 1,  
7.856030165834487: 1, 7.854266059964118: 1, 7.843427917138643: 1, 7.840871562  
158454: 1, 7.829259123537958: 1, 7.822818378790573: 1, 7.820793042708781: 1,  
7.820401656074233: 1, 7.820094868030294: 1, 7.812293830907494: 1, 7.799875600  
796433: 1, 7.7894243135781345: 1, 7.785869458882815: 1, 7.780305071986679: 1,  
7.77557314894385: 1, 7.775075380455164: 1, 7.768059638887945: 1, 7.7680427385  
14337: 1, 7.7530045163888985: 1, 7.749639116811307: 1, 7.738893191330873: 1,  
7.731819863607597: 1, 7.716676777107958: 1, 7.712596154745814: 1, 7.707583709  
461364: 1, 7.697043934266297: 1, 7.691110107328914: 1, 7.683740599918559: 1,  
7.678993481371007: 1, 7.676874435604698: 1, 7.675694163742105: 1, 7.672760653  
271402: 1, 7.665966035569618: 1, 7.656657352204757: 1, 7.65252263203452: 1,  
7.649732685565413: 1, 7.636024467193714: 1, 7.629917351312548: 1, 7.620955888  
071794: 1, 7.6015447463936: 1, 7.596486162014144: 1, 7.5923967682835345: 1,  
7.58549807879333: 1, 7.579979104607622: 1, 7.578546434116645: 1, 7.5765519885  
06318: 1, 7.5757148767474565: 1, 7.573513570668859: 1, 7.567729308850912: 1,  
7.567112828470178: 1, 7.565050322142518: 1, 7.561208835257127: 1, 7.552880012  
4009665: 1, 7.537706768010431: 1, 7.535302001347377: 1, 7.531305564176751: 1,  
7.52552794729805: 1, 7.524056530006512: 1, 7.5004912639074774: 1, 7.500368956  
5411795: 1, 7.494353748037146: 1, 7.493771133591274: 1, 7.492623575283266: 1,  
7.491914594380652: 1, 7.490465858403182: 1, 7.481605860894506: 1, 7.464852301  
690328: 1, 7.457897105867307: 1, 7.453723167873284: 1, 7.4338257822203975: 1,  
7.43205071248247: 1, 7.426882506110944: 1, 7.417724229312042: 1, 7.4156609775  
6492: 1, 7.407429070769639: 1, 7.407069248447494: 1, 7.395224756054181: 1, 7.

385027498918969: 1, 7.3819470280007735: 1, 7.376295363422325: 1, 7.3724160693  
59654: 1, 7.372168934412718: 1, 7.37056954351975: 1, 7.358697481142908: 1, 7.  
3567605805970775: 1, 7.34524697777243: 1, 7.331273447482039: 1, 7.33089710486  
8821: 1, 7.328834950109552: 1, 7.325348063947937: 1, 7.324513343679015: 1, 7.  
3231926599574075: 1, 7.311743325620549: 1, 7.302397362611505: 1, 7.2984641220  
27087: 1, 7.297428651857971: 1, 7.277788625821352: 1, 7.272069622891181: 1,  
7.2665605508413345: 1, 7.2661654985144235: 1, 7.262079633582819: 1, 7.2434359  
60393033: 1, 7.230768563286913: 1, 7.225318046248361: 1, 7.224601996062507:  
1, 7.222931691983325: 1, 7.208993595561583: 1, 7.203028236280691: 1, 7.198497  
556567738: 1, 7.189427289310809: 1, 7.188144049552843: 1, 7.185338183266622:  
1, 7.18529817691394: 1, 7.1827591387949274: 1, 7.176924726934586: 1, 7.169701  
16603311: 1, 7.168170338187: 1, 7.167533576490072: 1, 7.164591213711705: 1,  
7.162829408899581: 1, 7.150537313265291: 1, 7.145527576123244: 1, 7.139860601  
870029: 1, 7.125138716502345: 1, 7.1197326624881345: 1, 7.113935992326688: 1,  
7.086680818516647: 1, 7.078379824106145: 1, 7.076529204930513: 1, 7.074931343  
993786: 1, 7.074705862731387: 1, 7.068190397955633: 1, 7.062720125432995: 1,  
7.062642331534362: 1, 7.05886435676685: 1, 7.053358305585349: 1, 7.0523900065  
341945: 1, 7.051368268589938: 1, 7.046863796043614: 1, 7.038957163457341: 1,  
7.033446692354523: 1, 7.0210283715620765: 1, 7.019799909346458: 1, 7.01216829  
1908941: 1, 7.011694021876302: 1, 7.007719003018658: 1, 7.006023008567924: 1,  
6.9959401902388745: 1, 6.98468284329739: 1, 6.9805903252373245: 1, 6.97933445  
1339554: 1, 6.971098125877917: 1, 6.9554082288400725: 1, 6.954517803341356:  
1, 6.94392890762758: 1, 6.940093676373766: 1, 6.938466620167119: 1, 6.9347370  
483602555: 1, 6.9298749791095755: 1, 6.909704549436163: 1, 6.907845457791064:  
1, 6.901959484344427: 1, 6.898311104203807: 1, 6.894764557956705: 1, 6.894307  
824148758: 1, 6.887913513507078: 1, 6.885250412753548: 1, 6.878559909685794:  
1, 6.8735074198405375: 1, 6.871892362838734: 1, 6.854175974631295: 1, 6.85320  
8855738512: 1, 6.852695088904501: 1, 6.852479371045506: 1, 6.850529910927380  
5: 1, 6.848901019670414: 1, 6.847977676642771: 1, 6.847698981305847: 1, 6.841  
574098291193: 1, 6.834295879635943: 1, 6.831463850304009: 1, 6.82869023779485  
75: 1, 6.828175471884745: 1, 6.821282177221524: 1, 6.820347352376933: 1, 6.81  
6214959152492: 1, 6.810926934973605: 1, 6.803112095082661: 1, 6.8030972747034  
02: 1, 6.794529956854853: 1, 6.7903911116017435: 1, 6.78232376328574: 1, 6.77  
7583160605031: 1, 6.775915249339214: 1, 6.774530503634549: 1, 6.7735224203228  
91: 1, 6.769666568038315: 1, 6.764643238276599: 1, 6.751309937839854: 1, 6.74  
5173794883556: 1, 6.742013271024642: 1, 6.74093020068589: 1, 6.72985033008159  
6: 1, 6.726340394284921: 1, 6.699354858786156: 1, 6.698770112285369: 1, 6.689  
005374934036: 1, 6.68729382598378: 1, 6.684576240219602: 1, 6.68390670165736:  
1, 6.681762667767702: 1, 6.673987329714448: 1, 6.669151504768222: 1, 6.666958  
351142194: 1, 6.661267695476758: 1, 6.66124648141493: 1, 6.660812835520064:  
1, 6.650407003089762: 1, 6.647890434153005: 1, 6.64734754942142: 1, 6.6443683  
12676256: 1, 6.641620940430767: 1, 6.632954746664222: 1, 6.621418936891897:  
1, 6.620526460474487: 1, 6.612205716945513: 1, 6.61173678726514: 1, 6.5972884  
49353827: 1, 6.595988474359389: 1, 6.590359087750559: 1, 6.588224735913206:  
1, 6.587163473580585: 1, 6.577552745013945: 1, 6.576888470622014: 1, 6.569511  
4066511735: 1, 6.56718844356095: 1, 6.564037781687151: 1, 6.560358327961881:  
1, 6.552803869793501: 1, 6.549844933179496: 1, 6.549624354492555: 1, 6.541580  
297686458: 1, 6.535347223492063: 1, 6.53460794331718: 1, 6.532686506743115:  
1, 6.531507557993811: 1, 6.529916853096697: 1, 6.529279616995534: 1, 6.525091  
000664415: 1, 6.511909516968993: 1, 6.500566721564747: 1, 6.50044493232545:  
1, 6.493574685751356: 1, 6.492877804080287: 1, 6.49109404820549: 1, 6.4859416  
72926976: 1, 6.468077565327889: 1, 6.466009118249975: 1, 6.4657841287346205:  
1, 6.453741683621443: 1, 6.451064988245463: 1, 6.446786839376672: 1, 6.443740  
163383907: 1, 6.443328497894149: 1, 6.442063274999907: 1, 6.4322688890632165:  
1, 6.430388450024545: 1, 6.414174728703053: 1, 6.412904031538668: 1, 6.411988  
1816783995: 1, 6.409482842692491: 1, 6.406560257799249: 1, 6.398505130200458  
5: 1, 6.397573811915197: 1, 6.391072525463014: 1, 6.388345099858011: 1, 6.387



515652544683: 1, 6.37657002806621: 1, 6.375968690930346: 1, 6.37143688374543  
4: 1, 6.369663986862744: 1, 6.36427765591762: 1, 6.357485637959014: 1, 6.3558  
146841532155: 1, 6.353569166626952: 1, 6.350446525381036: 1, 6.34360864583478  
6: 1, 6.343515696610905: 1, 6.342846271921501: 1, 6.330114316805535: 1, 6.318  
510838774051: 1, 6.31076425869817: 1, 6.309377207073936: 1, 6.30783290334857  
7: 1, 6.300637022746812: 1, 6.3002771534142115: 1, 6.297436266528405: 1, 6.29  
7384824091763: 1, 6.296747839498341: 1, 6.293186404242105: 1, 6.2908182915794  
58: 1, 6.288218211761155: 1, 6.285955798483476: 1, 6.284656131636122: 1, 6.28  
4594835555546: 1, 6.278449528248002: 1, 6.277489834896366: 1, 6.2747265348684  
77: 1, 6.269938232948384: 1, 6.266859971245541: 1, 6.263543935437717: 1, 6.25  
5404014295078: 1, 6.240498955475491: 1, 6.2350549927997125: 1, 6.230784291799  
234: 1, 6.223093291494692: 1, 6.215807478288429: 1, 6.2151788412262645: 1, 6.  
210754851449287: 1, 6.207188072930951: 1, 6.198035626402154: 1, 6.18612763195  
2373: 1, 6.1767855743079325: 1, 6.173506632341786: 1, 6.171269851182426: 1,  
6.1710740393201515: 1, 6.169799839533524: 1, 6.155573851923314: 1, 6.14885926  
3596764: 1, 6.143640313807616: 1, 6.14312992322315: 1, 6.1320208784665695: 1,  
6.128467130836381: 1, 6.120705207179262: 1, 6.114179154816723: 1, 6.109773197  
1247455: 1, 6.103340374163565: 1, 6.102012947412282: 1, 6.090922503329984: 1,  
6.087949730722574: 1, 6.068762702964481: 1, 6.066073412103794: 1, 6.055142427  
860265: 1, 6.05315939727707: 1, 6.04898133909134: 1, 6.0465059003174435: 1,  
6.04408740068659: 1, 6.033234923616095: 1, 6.032738671834175: 1, 6.0326402380  
230135: 1, 6.031184360703106: 1, 6.030987914541547: 1, 6.027566910436541: 1,  
6.026188322247555: 1, 6.020403034597123: 1, 6.020384053285587: 1, 6.014339933  
300137: 1, 6.012582291479741: 1, 6.005418856249745: 1, 5.997918988459353: 1,  
5.995801759277665: 1, 5.9937885587966155: 1, 5.981621599436273: 1, 5.96330113  
3973628: 1, 5.963281235464185: 1, 5.957850874438121: 1, 5.955831943136923: 1,  
5.9510724189885815: 1, 5.940128428202666: 1, 5.93221525427665: 1, 5.931288025  
659112: 1, 5.916638237477272: 1, 5.905862372888252: 1, 5.9052447376936605: 1,  
5.903366434790971: 1, 5.902945796042184: 1, 5.898689280883562: 1, 5.894167980  
711571: 1, 5.892900188189126: 1, 5.892185013133433: 1, 5.888131343212597: 1,  
5.881935727955796: 1, 5.8796701738692265: 1, 5.879339079862091: 1, 5.87903135  
796333: 1, 5.878674793317592: 1, 5.873186054052349: 1, 5.872417267999007: 1,  
5.871716351528401: 1, 5.871545392294013: 1, 5.8689512786304: 1, 5.86724635678  
7675: 1, 5.866252754207397: 1, 5.865140510010267: 1, 5.864697550153495: 1, 5.  
863277516457579: 1, 5.827541000869115: 1, 5.820460543945737: 1, 5.82033843151  
8924: 1, 5.816609915855818: 1, 5.812473013029323: 1, 5.810921940532548: 1, 5.  
808729284544381: 1, 5.807362529109448: 1, 5.806761023161498: 1, 5.79948965658  
5405: 1, 5.795160049823191: 1, 5.794873374933489: 1, 5.791661119087586: 1, 5.  
789793164596802: 1, 5.786540646054724: 1, 5.7795738186452805: 1, 5.7780465826  
80203: 1, 5.777067251507555: 1, 5.776276099677748: 1, 5.764876577577711: 1,  
5.7628945516756875: 1, 5.7618090485124664: 1, 5.761457710635057: 1, 5.7534226  
95057535: 1, 5.746870806364432: 1, 5.7459469253980515: 1, 5.744517156137161:  
1, 5.735374105623882: 1, 5.734578031617064: 1, 5.731857274609318: 1, 5.724203  
101725122: 1, 5.714956597172376: 1, 5.705436191221437: 1, 5.698972084352073:  
1, 5.6970273401037606: 1, 5.696021917535018: 1, 5.688796139031355: 1, 5.68513  
5249899602: 1, 5.680966883163246: 1, 5.677249182544254: 1, 5.675707089953009  
5: 1, 5.675437518885716: 1, 5.6729280624336775: 1, 5.671534607999155: 1, 5.67  
1378729201887: 1, 5.667817324345495: 1, 5.666110706895406: 1, 5.6644599933718  
91: 1, 5.662361778936175: 1, 5.659338077513867: 1, 5.654949639289688: 1, 5.65  
1152739083522: 1, 5.64648153451257: 1, 5.6456408578683925: 1, 5.6455768785600  
9: 1, 5.642387656601188: 1, 5.63619658016733: 1, 5.635630846393857: 1, 5.6352  
21340681548: 1, 5.632936918052099: 1, 5.632661427044247: 1, 5.62721997978953  
5: 1, 5.622790841749682: 1, 5.6161493563275435: 1, 5.612023659247025: 1, 5.61  
0394600653861: 1, 5.609700829386011: 1, 5.6090255083173215: 1, 5.587642432069  
077: 1, 5.571540304302698: 1, 5.567584133209463: 1, 5.566548771163745: 1, 5.5  
64980508358852: 1, 5.560596427956347: 1, 5.559207369665225: 1, 5.556331395146  
322: 1, 5.554376597003859: 1, 5.550081273028058: 1, 5.547193415108779: 1, 5.5

44473416644668: 1, 5.544400878271412: 1, 5.539652790386198: 1, 5.539450604546  
589: 1, 5.53906053588975: 1, 5.538087433354151: 1, 5.537680979174492: 1, 5.53  
2216497137921: 1, 5.531759673215449: 1, 5.528803121680563: 1, 5.5277626039973  
34: 1, 5.527232643350295: 1, 5.525317823436987: 1, 5.522496152248665: 1, 5.51  
8442951850567: 1, 5.517501004378755: 1, 5.514789032844336: 1, 5.5131360501271  
96: 1, 5.509570762589798: 1, 5.509265929259796: 1, 5.502619362626053: 1, 5.49  
9502931485873: 1, 5.4978204026275455: 1, 5.487503438158661: 1, 5.487107851631  
773: 1, 5.480856450179883: 1, 5.478234173310868: 1, 5.475965353710287: 1, 5.4  
69877361339219: 1, 5.4654541369438485: 1, 5.459118148306031: 1, 5.45771256904  
6439: 1, 5.453377406042187: 1, 5.45265396168987: 1, 5.452589109415177: 1, 5.4  
4790746061192: 1, 5.437268425032982: 1, 5.43609944128947: 1, 5.43468453469821  
4: 1, 5.433649988177073: 1, 5.432823166833012: 1, 5.428995793962675: 1, 5.413  
745194587122: 1, 5.400927968954027: 1, 5.390493147575962: 1, 5.3887833288183  
9: 1, 5.387665791210405: 1, 5.386656895780422: 1, 5.384585840743472: 1, 5.379  
293969002137: 1, 5.379227250341898: 1, 5.377500046998303: 1, 5.37459629689579  
45: 1, 5.372585182011097: 1, 5.36171407454631: 1, 5.360579884511451: 1, 5.356  
999404743565: 1, 5.354303291990656: 1, 5.34736561514534: 1, 5.34491478550216  
7: 1, 5.336295728431571: 1, 5.335382053706493: 1, 5.3345367155924: 1, 5.33403  
1698366675: 1, 5.332838832753911: 1, 5.315871978330009: 1, 5.312350240322388:  
1, 5.312010451487504: 1, 5.303384849510536: 1, 5.299901531116548: 1, 5.297343  
954089837: 1, 5.295348187415909: 1, 5.294550623383171: 1, 5.281256613628355:  
1, 5.270653988803258: 1, 5.269199531168359: 1, 5.26789345517102: 1, 5.2673319  
2622349: 1, 5.2553448253517425: 1, 5.250369616477234: 1, 5.2412893680704915:  
1, 5.240603797175084: 1, 5.233052453649941: 1, 5.221022291338704: 1, 5.220154  
985900155: 1, 5.219897921881736: 1, 5.2162402136173975: 1, 5.214937517261312:  
1, 5.204758033966555: 1, 5.1891626176706716: 1, 5.180176874937202: 1, 5.17236  
5194737352: 1, 5.171020225891639: 1, 5.165988858022691: 1, 5.162722796942465:  
1, 5.159868194901314: 1, 5.157529211895454: 1, 5.1565084727591755: 1, 5.15532  
26885405135: 1, 5.151675186659846: 1, 5.146450657743205: 1, 5.145528617761840  
5: 1, 5.141127076156626: 1, 5.141083200602227: 1, 5.139781710817536: 1, 5.139  
688834414076: 1, 5.139089077388967: 1, 5.12076479255567: 1, 5.11966449958473  
1: 1, 5.117674234694734: 1, 5.113910118748087: 1, 5.112808779978584: 1, 5.107  
0355533490055: 1, 5.10651074469701: 1, 5.105062322330148: 1, 5.10453083101992  
7: 1, 5.103455441002889: 1, 5.101075269399237: 1, 5.089123378628886: 1, 5.081  
763476040172: 1, 5.079457295905482: 1, 5.066154778580886: 1, 5.06421850429222  
9: 1, 5.06366481361975: 1, 5.06154508401369: 1, 5.0613224879358425: 1, 5.0605  
84710646138: 1, 5.060365959813561: 1, 5.059024069555945: 1, 5.05877452699054  
6: 1, 5.058772407670942: 1, 5.056090353860968: 1, 5.054845978773364: 1, 5.054  
052081722371: 1, 5.053574214627069: 1, 5.049689577204261: 1, 5.04406760545964  
6: 1, 5.041720867945053: 1, 5.039799441326703: 1, 5.038179638534551: 1, 5.037  
468254064151: 1, 5.036286852397473: 1, 5.034551516911696: 1, 5.02984164572060  
5: 1, 5.027057478379977: 1, 5.025851467456503: 1, 5.022272845592937: 1, 5.021  
8516493604985: 1, 5.012002877192234: 1, 5.010956962729517: 1, 5.0107518102712  
45: 1, 5.00590894558583: 1, 4.9998222430877926: 1, 4.999406118296209: 1, 4.99  
5274238269705: 1, 4.9941250324549955: 1, 4.992969202769172: 1, 4.989890927562  
367: 1, 4.989126588358104: 1, 4.984636568144023: 1, 4.984079904521168: 1, 4.9  
83935748109963: 1, 4.978945608054387: 1, 4.975931593549244: 1, 4.968872880977  
529: 1, 4.9639864256453: 1, 4.95925440625599: 1, 4.957418109879136: 1, 4.9567  
30058278727: 1, 4.9542932862439955: 1, 4.953568555696982: 1, 4.95204608123638  
85: 1, 4.951984056520089: 1, 4.949504597714414: 1, 4.94910384983532: 1, 4.945  
831319069116: 1, 4.943708963757511: 1, 4.941152134276553: 1, 4.93923935140912  
1: 1, 4.9388249940126245: 1, 4.937250774777199: 1, 4.931420742741158: 1, 4.93  
0290984876541: 1, 4.929176597274912: 1, 4.928042710020301: 1, 4.9274257103494  
52: 1, 4.925761391897712: 1, 4.9255414855112765: 1, 4.92074932726269: 1, 4.90  
8602937587921: 1, 4.904722999980031: 1, 4.900898019101675: 1, 4.8962825397782  
22: 1, 4.894873208167134: 1, 4.893468547768292: 1, 4.893332566814321: 1, 4.89  
3173927356759: 1, 4.888128684295156: 1, 4.887885758590525: 1, 4.8712108072082

45: 1, 4.870150995738257: 1, 4.868437250277618: 1, 4.866216945042803: 1, 4.86  
1911814116857: 1, 4.850780277688517: 1, 4.846602844979642: 1, 4.8457685246290  
69: 1, 4.84110901210518: 1, 4.83842483970687: 1, 4.835153024863251: 1, 4.8319  
66902879494: 1, 4.831011547335977: 1, 4.830668597664459: 1, 4.82776079409738  
9: 1, 4.826534614513615: 1, 4.820418311378226: 1, 4.819486121638456: 1, 4.807  
738556319653: 1, 4.805542801357546: 1, 4.798551866414923: 1, 4.79767470990930  
3: 1, 4.794986809490969: 1, 4.794974892649809: 1, 4.7936351392475665: 1, 4.79  
35809865411665: 1, 4.790193386397884: 1, 4.786995777893809: 1, 4.784895085905  
547: 1, 4.78131600944506: 1, 4.78057216521009: 1, 4.77886162310863: 1, 4.7777  
30047984047: 1, 4.775928975997887: 1, 4.773960745590681: 1, 4.76989590968276  
3: 1, 4.768799069609183: 1, 4.768740861658539: 1, 4.763246649261218: 1, 4.758  
978427375315: 1, 4.758886718904766: 1, 4.758790172143131: 1, 4.75535007021487  
4: 1, 4.753758011774955: 1, 4.752901979869622: 1, 4.751985522616786: 1, 4.750  
864484318578: 1, 4.750441425506989: 1, 4.750025940768779: 1, 4.74955087741366  
1: 1, 4.747549134997196: 1, 4.745896619794648: 1, 4.74490589784158: 1, 4.7439  
96725907298: 1, 4.742068360598096: 1, 4.734088430709834: 1, 4.732371139486656  
6: 1, 4.724991393527957: 1, 4.722059816260948: 1, 4.719471377196999: 1, 4.719  
21056632071: 1, 4.716612576973192: 1, 4.7146343724807025: 1, 4.70780419457291  
1: 1, 4.700727612460848: 1, 4.699806380791066: 1, 4.698088117333121: 1, 4.697  
066267356256: 1, 4.694727580635782: 1, 4.6908608935358185: 1, 4.6889172214117  
05: 1, 4.684448165392596: 1, 4.683993547107189: 1, 4.683195100921038: 1, 4.67  
688630368336: 1, 4.676074727808154: 1, 4.6667690290744135: 1, 4.6660499745421  
44: 1, 4.6643728803387905: 1, 4.6636018650296975: 1, 4.662999024776231: 1, 4.  
660291378441485: 1, 4.656367404850576: 1, 4.655781775232488: 1, 4.65506222957  
4373: 1, 4.654190640633929: 1, 4.652876582467517: 1, 4.652251904177874: 1, 4.  
651815025177954: 1, 4.65049644745777: 1, 4.649656907377585: 1, 4.648692646634  
543: 1, 4.643686514270899: 1, 4.643426342905488: 1, 4.639180587445792: 1, 4.6  
36916423645925: 1, 4.635182232581991: 1, 4.634658667262613: 1, 4.632843204464  
626: 1, 4.630912328541166: 1, 4.627496862653686: 1, 4.6221014042579345: 1, 4.  
621160723504042: 1, 4.6175081829657785: 1, 4.615707565490562: 1, 4.6087905281  
81029: 1, 4.604406814830227: 1, 4.601368391728523: 1, 4.601264007948587: 1,  
4.594061836410239: 1, 4.593074808042741: 1, 4.592251279947166: 1, 4.592123187  
189329: 1, 4.58758786259062: 1, 4.5870436540547335: 1, 4.576212700212289: 1,  
4.574359636581015: 1, 4.573174071672549: 1, 4.572036921777054: 1, 4.562208607  
62768: 1, 4.560482067246486: 1, 4.559202617535509: 1, 4.555936003869667: 1,  
4.55316169195895: 1, 4.552843496455748: 1, 4.547750875700927: 1, 4.5446086750  
443415: 1, 4.54315633232467: 1, 4.543068596622529: 1, 4.53547706173976: 1, 4.  
531471694244074: 1, 4.5277120951760645: 1, 4.527248973111972: 1, 4.5271620984  
20333: 1, 4.527022629318957: 1, 4.526647729140743: 1, 4.52443748405229: 1, 4.  
516951871091057: 1, 4.516123744762207: 1, 4.515119274114703: 1, 4.51254632346  
7512: 1, 4.509080180114841: 1, 4.507898185017586: 1, 4.505865589680778: 1, 4.  
504159446291289: 1, 4.500856078366969: 1, 4.499623843406334: 1, 4.49929027328  
3084: 1, 4.49454630256541: 1, 4.4918585991235656: 1, 4.488065699819345: 1, 4.  
486953891225096: 1, 4.486870239712206: 1, 4.484912608124847: 1, 4.48464797173  
8969: 1, 4.48057148810334: 1, 4.478889244701549: 1, 4.478459049914942: 1, 4.4  
78106815507935: 1, 4.476808794056319: 1, 4.476785735302247: 1, 4.473496685495  
861: 1, 4.472834280704315: 1, 4.470158120372596: 1, 4.470098144695698: 1, 4.4  
68862361518404: 1, 4.46699466352917: 1, 4.464904814342533: 1, 4.4639567986005  
02: 1, 4.461882873957076: 1, 4.459943768152855: 1, 4.458357176196144: 1, 4.45  
6170081203698: 1, 4.456115829012881: 1, 4.4512874449838815: 1, 4.449297052006  
84: 1, 4.440534709574476: 1, 4.438105018468018: 1, 4.437542067664914: 1, 4.43  
4013330283239: 1, 4.4298895375527065: 1, 4.4297671091770106: 1, 4.42931186841  
8367: 1, 4.422440788205055: 1, 4.42127309269398: 1, 4.420917140346345: 1, 4.4  
20784551633854: 1, 4.418835131408685: 1, 4.414328514727713: 1, 4.413474450207  
572: 1, 4.413229130028185: 1, 4.4117481879257605: 1, 4.411440727279674: 1, 4.  
405400760570337: 1, 4.40240933436706: 1, 4.3997023058762466: 1, 4.39774149355  
1306: 1, 4.39669283557897: 1, 4.394956171498758: 1, 4.389043065971306: 1, 4.3

87704972370426: 1, 4.38263125453914: 1, 4.382300204154238: 1, 4.3767700689760  
28: 1, 4.366082004734845: 1, 4.365421938280668: 1, 4.360675151832565: 1, 4.35  
6627292289897: 1, 4.353017732674497: 1, 4.348551223841266: 1, 4.3478644931977  
16: 1, 4.346360719468213: 1, 4.346109263623332: 1, 4.341347354677878: 1, 4.33  
5766451094374: 1, 4.332849634571223: 1, 4.331979591661529: 1, 4.3301894343662  
36: 1, 4.3296647158202095: 1, 4.329034667127301: 1, 4.3267978450449975: 1, 4.  
321467903974046: 1, 4.319409168910403: 1, 4.318389284653249: 1, 4.31749397988  
6841: 1, 4.316283809859334: 1, 4.308700033575829: 1, 4.308037052873442: 1, 4.  
30485222254125: 1, 4.302819548758459: 1, 4.301785101886079: 1, 4.300221573735  
576: 1, 4.293850303824495: 1, 4.2925658674443845: 1, 4.292339294438951: 1, 4.  
291570801294596: 1, 4.2876353316794935: 1, 4.28527299690806: 1, 4.28318845140  
7836: 1, 4.279475280185159: 1, 4.279216676365495: 1, 4.279084020979221: 1, 4.  
275346166749218: 1, 4.274419760930619: 1, 4.272218921615013: 1, 4.26845192034  
0591: 1, 4.267009818898967: 1, 4.26294783156286: 1, 4.2608420147522565: 1, 4.  
260763386165405: 1, 4.257160230396238: 1, 4.256162400088198: 1, 4.25613768875  
3975: 1, 4.255156647041396: 1, 4.253393731997442: 1, 4.25178998767272: 1, 4.2  
443107941057505: 1, 4.242404599181007: 1, 4.238525534887753: 1, 4.23630533566  
7456: 1, 4.235355845532031: 1, 4.234188725864385: 1, 4.233541352537479: 1, 4.  
23330156640022: 1, 4.232084744584502: 1, 4.228357728312868: 1, 4.227266322522  
207: 1, 4.22724426636135: 1, 4.226979108590222: 1, 4.226732186326534: 1, 4.22  
512759170856: 1, 4.221147685309821: 1, 4.220940744278355: 1, 4.22082177744649  
8: 1, 4.218819284022429: 1, 4.215956541517889: 1, 4.215618223236915: 1, 4.214  
217204720615: 1, 4.211668757205356: 1, 4.209612951289684: 1, 4.20861585924912  
25: 1, 4.205561825133731: 1, 4.203162848826229: 1, 4.1951476837506965: 1, 4.1  
93202822274016: 1, 4.192027825142278: 1, 4.190787153581125: 1, 4.190690570675  
403: 1, 4.18931420419578: 1, 4.186565690829107: 1, 4.181475371867971: 1, 4.17  
7975857459942: 1, 4.177061633744532: 1, 4.174641363358141: 1, 4.1733012815772  
71: 1, 4.172670748374196: 1, 4.170462082541809: 1, 4.170367835352791: 1, 4.16  
9874778813473: 1, 4.166415934837324: 1, 4.166330700489613: 1, 4.1657585657039  
17: 1, 4.165687832175082: 1, 4.165405097721872: 1, 4.158923338119645: 1, 4.15  
6319457839494: 1, 4.1542752967030605: 1, 4.151509729894343: 1, 4.149644594371  
968: 1, 4.14561420215654: 1, 4.14516541260324: 1, 4.144713456168646: 1, 4.142  
440843057854: 1, 4.139506089918387: 1, 4.138748357228107: 1, 4.13755882273408  
7: 1, 4.137141815395755: 1, 4.134476123275039: 1, 4.133902702786823: 1, 4.128  
067921256107: 1, 4.126602104319434: 1, 4.121795655327156: 1, 4.11690163616694  
8: 1, 4.115425590303513: 1, 4.111319784162429: 1, 4.110509890372155: 1, 4.108  
932919916527: 1, 4.108387113275039: 1, 4.105087599662516: 1, 4.10256462300490  
8: 1, 4.101307691591375: 1, 4.095273026614111: 1, 4.094620214327885: 1, 4.092  
8492142561765: 1, 4.091819273459392: 1, 4.0906413456693205: 1, 4.087571180708  
53: 1, 4.086319359590117: 1, 4.085702053218294: 1, 4.0854533822745305: 1, 4.0  
85325849980406: 1, 4.084741567849183: 1, 4.082160997998085: 1, 4.079610998470  
426: 1, 4.079186251231372: 1, 4.078300367972983: 1, 4.07777853668905: 1, 4.07  
2894846073083: 1, 4.070338479334589: 1, 4.066972117254546: 1, 4.0667504517610  
7: 1, 4.065894750531521: 1, 4.064296593474717: 1, 4.061835600203901: 1, 4.059  
175516635944: 1, 4.058938124695932: 1, 4.058778241223024: 1, 4.0587038009675  
3: 1, 4.056378358539442: 1, 4.049106557974373: 1, 4.048818506140668: 1, 4.045  
014097095545: 1, 4.043669631632621: 1, 4.041979436744445: 1, 4.04179135444109  
5: 1, 4.041105374179748: 1, 4.039879378016334: 1, 4.038183011448725: 1, 4.036  
672439258585: 1, 4.036055401154339: 1, 4.034947990226329: 1, 4.03080333347998  
6: 1, 4.027791128669757: 1, 4.026732568995539: 1, 4.026375227116831: 1, 4.024  
809351312656: 1, 4.022176005415525: 1, 4.0210720969896725: 1, 4.0206659800196  
2: 1, 4.020553995109056: 1, 4.019661674496704: 1, 4.018097874673741: 1, 4.016  
229737974829: 1, 4.015849822482197: 1, 4.010143258774422: 1, 4.00693526528394  
1: 1, 4.00610676303657: 1, 4.006023308266325: 1, 3.9987007842420352: 1, 3.998  
665326977374: 1, 3.998378487620518: 1, 3.9945123627939636: 1, 3.9940515992091  
26: 1, 3.9897718375446436: 1, 3.9861722587796753: 1, 3.9823933813790235: 1,  
3.9822030588921793: 1, 3.981248161101295: 1, 3.9803139622435335: 1, 3.9781889

458755764: 1, 3.974626788768033: 1, 3.9737087129588406: 1, 3.969381897149635  
3: 1, 3.969307084738411: 1, 3.9690265069745885: 1, 3.962771060457636: 1, 3.95  
9603467363326: 1, 3.955868226887048: 1, 3.955470561614359: 1, 3.9553612900009  
987: 1, 3.9546163345605807: 1, 3.951690626787508: 1, 3.951278626975126: 1, 3.  
950908025785735: 1, 3.950845463996645: 1, 3.94785250797679: 1, 3.942357747944  
2384: 1, 3.942170081223109: 1, 3.9393735861447237: 1, 3.9379178033391815: 1,  
3.934756826586977: 1, 3.9344863262255254: 1, 3.93274527610205: 1, 3.931642288  
456989: 1, 3.9261135020132896: 1, 3.924246222130553: 1, 3.923122107674109: 1,  
3.920811708082094: 1, 3.9200523730034926: 1, 3.9171139639454555: 1, 3.9170374  
73604068: 1, 3.9158561037546784: 1, 3.9140477481925013: 1, 3.913776528476961  
7: 1, 3.906821448215659: 1, 3.905090363565722: 1, 3.9050543648745473: 1, 3.90  
16863446754155: 1, 3.9015440772585466: 1, 3.892909862562512: 1, 3.88696612747  
78498: 1, 3.8864923132062934: 1, 3.8836653536696852: 1, 3.8806198451374194:  
1, 3.8746322050372575: 1, 3.8732131645955987: 1, 3.8702231136841774: 1, 3.866  
0877286314608: 1, 3.8658990821640282: 1, 3.8619297587103656: 1, 3.86035256581  
6864: 1, 3.858695535940779: 1, 3.8572802107730233: 1, 3.8549262410964436: 1,  
3.8544422991361436: 1, 3.8533016794889616: 1, 3.8516475651791966: 1, 3.846369  
1682670818: 1, 3.8440479389994984: 1, 3.8433829731437315: 1, 3.83758718446652  
44: 1, 3.835634719722884: 1, 3.834620437036863: 1, 3.833861282154509: 1, 3.82  
72144320517754: 1, 3.824640087083814: 1, 3.8228754600800077: 1, 3.82279253233  
40517: 1, 3.822786751629589: 1, 3.82010262748874: 1, 3.8192593688067826: 1,  
3.8137546719821627: 1, 3.8075480338421035: 1, 3.8030597813462905: 1, 3.802861  
4638710287: 1, 3.802471866960158: 1, 3.802232258086435: 1, 3.800599806283066  
6: 1, 3.7987289227494303: 1, 3.7959443039111305: 1, 3.794749625631295: 1, 3.7  
93791768973213: 1, 3.793019551091195: 1, 3.792350406136747: 1, 3.786984904908  
487: 1, 3.7844344708648925: 1, 3.78221680717954: 1, 3.7822131590657784: 1, 3.  
781100199072917: 1, 3.7787532688050276: 1, 3.77443921189116: 1, 3.76841106663  
49504: 1, 3.7650550205476407: 1, 3.7621844919247045: 1, 3.761428759323274: 1,  
3.7593223155440416: 1, 3.7592211555885755: 1, 3.756624481643063: 1, 3.7552341  
861026393: 1, 3.754974867491729: 1, 3.751385597611516: 1, 3.751194939842891:  
1, 3.7511279983942076: 1, 3.7510651113376463: 1, 3.749979784806412: 1, 3.7498  
93719127949: 1, 3.7438231913071784: 1, 3.7415032289517405: 1, 3.7412646410862  
602: 1, 3.7385938659455347: 1, 3.737237164824365: 1, 3.7298330636326966: 1,  
3.7288825362088236: 1, 3.7284219703185957: 1, 3.7282575617374083: 1, 3.727870  
5847669715: 1, 3.723838719303063: 1, 3.719892835400177: 1, 3.718673360372446  
4: 1, 3.7186628238987915: 1, 3.7180883810098804: 1, 3.715854128035774: 1, 3.7  
156382382657047: 1, 3.7148105758703176: 1, 3.713491022720891: 1, 3.7123151761  
93821: 1, 3.710608092255818: 1, 3.7099118386915224: 1, 3.7072075560610065: 1,  
3.7049101568287597: 1, 3.703993778548759: 1, 3.7000253654148505: 1, 3.6981730  
783346642: 1, 3.697053299150337: 1, 3.694922035162427: 1, 3.694619106570001:  
1, 3.6930326032449305: 1, 3.691039277524557: 1, 3.6888382753627016: 1, 3.6857  
94057393566: 1, 3.685049159629853: 1, 3.683287401896022: 1, 3.68171894951416  
5: 1, 3.681313572711852: 1, 3.6800548694235062: 1, 3.6800249178898783: 1, 3.6  
79634745347264: 1, 3.6788817217545153: 1, 3.677111096683095: 1, 3.67709762052  
6659: 1, 3.674905090452072: 1, 3.6748853686410192: 1, 3.66685773982135: 1, 3.  
664603269763988: 1, 3.6641960432522467: 1, 3.6640988504836445: 1, 3.662906074  
3784447: 1, 3.6627692058295875: 1, 3.6619511465681764: 1, 3.6584304882911813:  
1, 3.658297306895233: 1, 3.656346302290703: 1, 3.6563234487907543: 1, 3.64865  
93153578566: 1, 3.6482824019107993: 1, 3.645649471748178: 1, 3.64365659653065  
24: 1, 3.642148301965277: 1, 3.636865852602486: 1, 3.6360703320594245: 1, 3.6  
31748532694951: 1, 3.62756017277359: 1, 3.6257611164535795: 1, 3.618764902768  
234: 1, 3.616810301614538: 1, 3.6159203153988257: 1, 3.615625615120139: 1, 3.  
615373568109403: 1, 3.611348414525137: 1, 3.611274163928122: 1, 3.61120343462  
5296: 1, 3.607244837577702: 1, 3.6062470592530858: 1, 3.6046822588488427: 1,  
3.6040013467472267: 1, 3.6035111130403283: 1, 3.601480884363173: 1, 3.6014258  
651632427: 1, 3.6013974650530396: 1, 3.598818703540057: 1, 3.597842198932131:  
1, 3.5961602782886892: 1, 3.595195565312551: 1, 3.5924430720756644: 1, 3.5904

978413737982: 1, 3.5882490931514623: 1, 3.587540302305619: 1, 3.5867669583594  
22: 1, 3.5785722117992775: 1, 3.5748234906154446: 1, 3.574549743519945: 1, 3.  
573981561308301: 1, 3.566806955058511: 1, 3.56657278749464: 1, 3.564607057726  
1537: 1, 3.5637681324675454: 1, 3.5635906550320824: 1, 3.558260635651991: 1,  
3.5552162218186796: 1, 3.5550485437207784: 1, 3.5523784134225243: 1, 3.548205  
1108618826: 1, 3.546271376024326: 1, 3.5435464141329747: 1, 3.542001078047879  
7: 1, 3.5409409144296204: 1, 3.53965022679303: 1, 3.5380270308725486: 1, 3.53  
7341446986883: 1, 3.5371714281496107: 1, 3.536506542575079: 1, 3.535586263707  
5353: 1, 3.5333993517714557: 1, 3.5322592063369065: 1, 3.53066254471795: 1,  
3.5270426895167093: 1, 3.5265001715073256: 1, 3.5260889014005636: 1, 3.525843  
195811141: 1, 3.525164878816009: 1, 3.52293754696992: 1, 3.520131346547001:  
1, 3.519429992752532: 1, 3.517319293883311: 1, 3.515463218414912: 1, 3.510918  
070530162: 1, 3.5093687239798745: 1, 3.5092588392161637: 1, 3.508070382646068  
6: 1, 3.5080038711912773: 1, 3.5073427544141547: 1, 3.5060896424888393: 1, 3.  
5060795726731975: 1, 3.5046347220285066: 1, 3.502704578884599: 1, 3.501774172  
1414: 1, 3.501105697740837: 1, 3.4978653775130883: 1, 3.4948083062013318: 1,  
3.4933949499040575: 1, 3.4909970831677177: 1, 3.4902785549092608: 1, 3.482960  
3502360413: 1, 3.480239620742044: 1, 3.4793435429979946: 1, 3.476611361674562  
4: 1, 3.4757860457762058: 1, 3.4718375440212115: 1, 3.471790205405706: 1, 3.4  
70894582132835: 1, 3.4702591778257523: 1, 3.4667315276404396: 1, 3.4649434686  
29618: 1, 3.463642733578678: 1, 3.4619738592275118: 1, 3.4618035134360174: 1,  
3.4606419675673585: 1, 3.460254721741273: 1, 3.46012776233547: 1, 3.459339726  
9886843: 1, 3.4591245717558734: 1, 3.455770431742814: 1, 3.4539183094621144:  
1, 3.4528635569029893: 1, 3.4527142711182415: 1, 3.4526908200942783: 1, 3.452  
2030365181853: 1, 3.4519081657190878: 1, 3.449842380435351: 1, 3.449356329458  
241: 1, 3.448411596789727: 1, 3.4465381872371696: 1, 3.4440826856852973: 1,  
3.441762753257424: 1, 3.44058530112013: 1, 3.440270077925744: 1, 3.4393745187  
245477: 1, 3.4387245722019086: 1, 3.4369173334044034: 1, 3.432662774363207:  
1, 3.4320426023294424: 1, 3.430462570333492: 1, 3.4293700133203235: 1, 3.4256  
63733216876: 1, 3.4252286430547336: 1, 3.4224630256346007: 1, 3.4221574023273  
71: 1, 3.4221529479061616: 1, 3.4197365087409826: 1, 3.415870533040233: 1, 3.  
414282261552821: 1, 3.4104146581530017: 1, 3.405703004781899: 1, 3.4037731950  
209964: 1, 3.40359655451417: 1, 3.401226511471915: 1, 3.39902482452444: 1, 3.  
3988975747268997: 1, 3.3980275106056936: 1, 3.396993587833546: 1, 3.396877720  
1468754: 1, 3.3945275193858673: 1, 3.3942172906911914: 1, 3.3922947548565436:  
1, 3.39058287500259: 1, 3.3891589384555036: 1, 3.3875877329980506: 1, 3.38690  
7731657596: 1, 3.3866477705779348: 1, 3.3828350752399716: 1, 3.38170325528811  
37: 1, 3.373495349584724: 1, 3.37324684415292: 1, 3.3732407555247743: 1, 3.37  
2058998878573: 1, 3.3719970485782866: 1, 3.369950955940868: 1, 3.367568342843  
2266: 1, 3.3647994448180816: 1, 3.3596823767237898: 1, 3.358586662710163: 1,  
3.3565814025503564: 1, 3.351620828116798: 1, 3.3513409118734696: 1, 3.3500771  
12705513: 1, 3.3495382909778373: 1, 3.349142694059623: 1, 3.3476965352316714:  
1, 3.346961104314669: 1, 3.3467407335425072: 1, 3.346450202086028: 1, 3.34580  
40224526164: 1, 3.3453852659603673: 1, 3.343134429298677: 1, 3.33862766393500  
54: 1, 3.332087924203746: 1, 3.3310843869989895: 1, 3.323633333089102: 1, 3.3  
235238521449357: 1, 3.32321034686937: 1, 3.3213846511052263: 1, 3.32072536610  
1866: 1, 3.320568292026711: 1, 3.3190617657557113: 1, 3.315703052737846: 1,  
3.315502991656965: 1, 3.3092793923428148: 1, 3.309266203434792: 1, 3.30916323  
38507884: 1, 3.30639698564936: 1, 3.3060559781074925: 1, 3.306047519197358:  
1, 3.303657584502134: 1, 3.3034927922249553: 1, 3.30197571565142: 1, 3.301393  
3151391237: 1, 3.299856397700352: 1, 3.299304236365207: 1, 3.299300110738183  
6: 1, 3.295103179665053: 1, 3.2915734799888416: 1, 3.2903670072063447: 1, 3.2  
90051887185097: 1, 3.2867795763039482: 1, 3.286397208364785: 1, 3.27944840071  
79503: 1, 3.275214248313328: 1, 3.2747632596180534: 1, 3.269944199710947: 1,  
3.2687728117993133: 1, 3.2638490926282513: 1, 3.2638269192386677: 1, 3.263538  
5346271706: 1, 3.2611501971839227: 1, 3.2602973955609302: 1, 3.25906257137900  
33: 1, 3.2573763404386415: 1, 3.256646541586142: 1, 3.2522731068545827: 1, 3.

2507603713237736: 1, 3.243020054976453: 1, 3.240734765327281: 1, 3.2391135936  
154773: 1, 3.2383368086692053: 1, 3.237471768030247: 1, 3.2355740586182966:  
1, 3.233314361396216: 1, 3.2318921108416205: 1, 3.2294843076922146: 1, 3.2285  
186266247097: 1, 3.227979718072568: 1, 3.2253845275730826: 1, 3.2242590306043  
47: 1, 3.216792511072312: 1, 3.2147126234208643: 1, 3.213662817462621: 1, 3.2  
127657158047422: 1, 3.211268212393672: 1, 3.2090930373130075: 1, 3.2085941427  
810605: 1, 3.2080512192852946: 1, 3.207215991122145: 1, 3.206586798234558: 1,  
3.2059643204853043: 1, 3.2050032251655414: 1, 3.203408983002889: 1, 3.1989615  
20591315: 1, 3.1963479243739643: 1, 3.1931675853452273: 1, 3.192299991472078  
4: 1, 3.189767386312025: 1, 3.189667334430986: 1, 3.188552032389875: 1, 3.188  
0516060947697: 1, 3.1872857610766836: 1, 3.1864831392733426: 1, 3.18425224602  
95286: 1, 3.184052109437738: 1, 3.1835035393886355: 1, 3.1828680661014617: 1,  
3.1816917273002647: 1, 3.1804721454306417: 1, 3.1802678288712807: 1, 3.179900  
6245043033: 1, 3.179073597931619: 1, 3.176532729088297: 1, 3.173326431462191:  
1, 3.170679392660777: 1, 3.169013423662789: 1, 3.164920433924728: 1, 3.163545  
4658074917: 1, 3.1629355462682907: 1, 3.1620816244925742: 1, 3.15728265280414  
58: 1, 3.1569599397642465: 1, 3.154743475517321: 1, 3.152004929838707: 1, 3.1  
510991409010267: 1, 3.1502410672862373: 1, 3.1487371492956227: 1, 3.148689759  
1316305: 1, 3.1443523321457745: 1, 3.1414157744237503: 1, 3.1394780096463424:  
1, 3.136691169664281: 1, 3.1363047747382944: 1, 3.135121408405816: 1, 3.13398  
19999542953: 1, 3.131853212956354: 1, 3.1291437029519176: 1, 3.12759709242146  
84: 1, 3.1248026907705397: 1, 3.1227417473620416: 1, 3.119700548594342: 1, 3.  
1186714576268804: 1, 3.1176475535371835: 1, 3.1171322313934744: 1, 3.11584098  
49252453: 1, 3.1151610936404412: 1, 3.1142465033681845: 1, 3.111174298276123  
5: 1, 3.1105385796201035: 1, 3.1100731245608664: 1, 3.1067896546208336: 1, 3.  
104337483732121: 1, 3.1017019810570914: 1, 3.1015910162799063: 1, 3.100608512  
8497247: 1, 3.1003536612471883: 1, 3.0987924868763934: 1, 3.097270577976514:  
1, 3.09593039215193: 1, 3.0909883782868217: 1, 3.0893030476905246: 1, 3.08913  
9892555264: 1, 3.0877044672894947: 1, 3.086060974290635: 1, 3.085807163403102  
6: 1, 3.0839315468456863: 1, 3.0837294233164627: 1, 3.0787521417707056: 1, 3.  
0782509214485416: 1, 3.077482961310642: 1, 3.073906927194865: 1, 3.0725419157  
66331: 1, 3.072292595705374: 1, 3.069360016391759: 1, 3.0667313426104252: 1,  
3.0632518711521466: 1, 3.0593204881722893: 1, 3.057398530952036: 1, 3.0520869  
723216477: 1, 3.05197091865621: 1, 3.0501173725195283: 1, 3.049696350721832:  
1, 3.0475662398060632: 1, 3.046898324091203: 1, 3.0460252521651094: 1, 3.0440  
462940619817: 1, 3.0427411348364886: 1, 3.04077763417554: 1, 3.03937384913007  
7: 1, 3.0360643455846907: 1, 3.035743479814871: 1, 3.034868811964755: 1, 3.03  
26757266970796: 1, 3.0320885578948924: 1, 3.030294370142614: 1, 3.02955842413  
89366: 1, 3.028708748088951: 1, 3.0280846193206874: 1, 3.027854941806662: 1,  
3.0276423551952156: 1, 3.0229324180048094: 1, 3.0164615115987226: 1, 3.015366  
907056022: 1, 3.014286198598259: 1, 3.013666567042025: 1, 3.009999174167348:  
1, 3.009120249664613: 1, 3.0077269495772057: 1, 3.007484415219563: 1, 3.00734  
8730377347: 1, 3.0060285344585482: 1, 3.00398453448953: 1, 2.998846169040937  
6: 1, 2.9967841395238737: 1, 2.9949407709076286: 1, 2.993243266674515: 1, 2.9  
928309639430597: 1, 2.9902407391761185: 1, 2.9901229381438212: 1, 2.990075609  
3914715: 1, 2.989803244022582: 1, 2.9888377424417962: 1, 2.987359675668571:  
1, 2.985563687689467: 1, 2.9853828335064856: 1, 2.9845293849726278: 1, 2.9819  
883615744462: 1, 2.9818193892465716: 1, 2.9802611193921367: 1, 2.974241690663  
377: 1, 2.971353621893581: 1, 2.9712589632673425: 1, 2.970634425289743: 1, 2.  
969415785433234: 1, 2.969220774662364: 1, 2.968670124120835: 1, 2.96786391120  
1126: 1, 2.9668389736819205: 1, 2.9662439774523044: 1, 2.9644671243994813: 1,  
2.962615481804976: 1, 2.9603704777037447: 1, 2.960109499316348: 1, 2.95849705  
07691564: 1, 2.957889551506404: 1, 2.9551844330759205: 1, 2.954685546459267:  
1, 2.954410507320464: 1, 2.9535105571890297: 1, 2.9519852254022387: 1, 2.9505  
793883585456: 1, 2.950144431468807: 1, 2.945737511174325: 1, 2.9455982854048  
8: 1, 2.942084524275321: 1, 2.940854105999663: 1, 2.9399708140223444: 1, 2.93  
9086092466388: 1, 2.937596856042709: 1, 2.937094135956869: 1, 2.9362597556463

057: 1, 2.9340767379194017: 1, 2.9340400450513715: 1, 2.9336669324466684: 1, 2.933371989026639: 1, 2.9331882083587386: 1, 2.9329359830233734: 1, 2.9323977968374146: 1, 2.931652853960488: 1, 2.926611526819323: 1, 2.9256798321314683: 1, 2.925229060566934: 1, 2.924992110587143: 1, 2.9233260198437643: 1, 2.923014243373859: 1, 2.9215383368794114: 1, 2.92126849632764: 1, 2.9207373791891875: 1, 2.916057717489083: 1, 2.914394728520284: 1, 2.9142742024492976: 1, 2.9076236101529926: 1, 2.9074907354412503: 1, 2.9060266156040337: 1, 2.904831924452731: 1, 2.904479564102875: 1, 2.8971678378666406: 1, 2.891779620550435: 1, 2.8891645728403397: 1, 2.8885170357466237: 1, 2.8882358490562843: 1, 2.8876391291882815: 1, 2.887588254081636: 1, 2.887202003495418: 1, 2.886032628054872: 1, 2.8812826021763205: 1, 2.8796205020116643: 1, 2.8792937678873596: 1, 2.8786892502153694: 1, 2.878032990079043: 1, 2.877425878030673: 1, 2.8759452272966506: 1, 2.873129728663176: 1, 2.8730974153573223: 1, 2.8721550285526427: 1, 2.871637182011215: 1, 2.8669553862564556: 1, 2.863655249226294: 1, 2.862263783558363: 1, 2.860122769198284: 1, 2.8593080783421274: 1, 2.8579308711969786: 1, 2.8558291669680007: 1, 2.855256462875683: 1, 2.8514989065378358: 1, 2.8478507696449498: 1, 2.8431311689130854: 1, 2.840793028975837: 1, 2.8406609899479: 1, 2.839894338767075: 1, 2.839713299437253: 1, 2.8397101567584184: 1, 2.8396602795763326: 1, 2.8388109785354434: 1, 2.838235007862375: 1, 2.837933829326647: 1, 2.830115928609448: 1, 2.8266347945352845: 1, 2.8262330420571335: 1, 2.825878751665906: 1, 2.8236373434968467: 1, 2.823524417246526: 1, 2.8226238503450456: 1, 2.8205624327222933: 1, 2.8202381858653665: 1, 2.819468312330465: 1, 2.819043574651522: 1, 2.8150923410319426: 1, 2.813796649047008: 1, 2.8134194787086346: 1, 2.811162555721237: 1, 2.808793065755371: 1, 2.807703400857976: 1, 2.8071535259034404: 1, 2.8060723521443474: 1, 2.7999805530822006: 1, 2.7987025090810054: 1, 2.7981157009388444: 1, 2.797210415719662: 1, 2.7951698078386866: 1, 2.7941012761846986: 1, 2.7917534336574494: 1, 2.7912548155414076: 1, 2.7888242495084863: 1, 2.787843997721283: 1, 2.7871166061864545: 1, 2.7857725382026874: 1, 2.7848789479932283: 1, 2.7812244567029794: 1, 2.7806644944343764: 1, 2.778047628341239: 1, 2.7766288453816053: 1, 2.7761794202705037: 1, 2.7726553974790447: 1, 2.7717995670736735: 1, 2.771429100173963: 1, 2.770492984678984: 1, 2.7703667155453893: 1, 2.7695679108325124: 1, 2.7695619068912984: 1, 2.7666134894253367: 1, 2.764575755741644: 1, 2.7644530550949593: 1, 2.7626918843980834: 1, 2.761947536475873: 1, 2.759395540249808: 1, 2.7592564122450236: 1, 2.7579436014644707: 1, 2.757624088959089: 1, 2.757447935768877: 1, 2.757111608656866: 1, 2.756130048338187: 1, 2.7534509584158964: 1, 2.7521484578043434: 1, 2.751645223298127: 1, 2.748828033304803: 1, 2.748583203077956: 1, 2.7481943568838996: 1, 2.7448091478331054: 1, 2.7433357864839856: 1, 2.741089604759948: 1, 2.740031821796593: 1, 2.7394401808783337: 1, 2.7390482902859548: 1, 2.736965477267326: 1, 2.736393984886058: 1, 2.735691748858124: 1, 2.7346702381155037: 1, 2.734021825761068: 1, 2.7338101782841733: 1, 2.7305763889767856: 1, 2.725034564191521: 1, 2.7237962959189286: 1, 2.7202601651538307: 1, 2.719219281994074: 1, 2.7171033168600873: 1, 2.716947985656484: 1, 2.7166749489991915: 1, 2.7163723877992996: 1, 2.715513961698774: 1, 2.715061758738887: 1, 2.7150196168217255: 1, 2.714087501897477: 1, 2.712219172816906: 1, 2.7116374328844417: 1, 2.7098250471929433: 1, 2.70819069647033: 1, 2.707090362818828: 1, 2.7066069149971277: 1, 2.7055481393363383: 1, 2.70511035825843: 1, 2.703529370065931: 1, 2.700786819019665: 1, 2.699898076421519: 1, 2.694420388027162: 1, 2.6918985301091807: 1, 2.685471634546622: 1, 2.684795249205513: 1, 2.679350325213471: 1, 2.6782739842637837: 1, 2.675438175373467: 1, 2.674230875595458: 1, 2.667857434934005: 1, 2.6677807155363222: 1, 2.6676550575857045: 1, 2.667479907645604: 1, 2.6660449170571545: 1, 2.665480729775682: 1, 2.6648849675525614: 1, 2.6648437784853916: 1, 2.6647042949285167: 1, 2.661771256685744: 1, 2.657752025130667: 1, 2.656794059336451: 1, 2.652852739204001: 1, 2.6519645868369506: 1, 2.6511742770286952: 1, 2.6476437429888606: 1, 2.647641775400457: 1, 2.6432950451536654: 1, 2.641351778067206: 1, 2.641181308435179: 1, 2.6408419894231803: 1, 2.640239968097961: 1, 2.6398991096216684: 1, 2.6380850



276009222: 1, 2.634726441002518: 1, 2.6333490759375744: 1, 2.632261013135168  
7: 1, 2.62807746666392: 1, 2.6217810637728505: 1, 2.6196337045580815: 1, 2.61  
9617575837605: 1, 2.6194904417981757: 1, 2.619036371500831: 1, 2.618273009408  
3984: 1, 2.6172056191878483: 1, 2.6168981530950117: 1, 2.6163427416502154: 1,  
2.615657171991476: 1, 2.6141483772608796: 1, 2.613926669049378: 1, 2.61290500  
36489843: 1, 2.6100086055300618: 1, 2.6091774261283454: 1, 2.607996863194966:  
1, 2.6062857618330626: 1, 2.6044697530456102: 1, 2.603827337682411: 1, 2.6012  
56110565748: 1, 2.601012055695202: 1, 2.600132133203487: 1, 2.598126585706028  
3: 1, 2.5969565661337133: 1, 2.595340556665602: 1, 2.59279450374548: 1, 2.590  
61472846246: 1, 2.587623026285294: 1, 2.587558293140676: 1, 2.58741374011548  
8: 1, 2.5823122028751335: 1, 2.5819591033885527: 1, 2.57896002549166: 1, 2.57  
79820230632433: 1, 2.5768772115689527: 1, 2.5750556020964384: 1, 2.5718007650  
83635: 1, 2.5697435444084995: 1, 2.5688115696942515: 1, 2.5683398441967324:  
1, 2.567808001200652: 1, 2.5664364914159155: 1, 2.560902319207104: 1, 2.56077  
70797098337: 1, 2.55946381498254: 1, 2.5593697509966264: 1, 2.557862565665057  
5: 1, 2.557095252881177: 1, 2.5565660091899662: 1, 2.5548550894642017: 1, 2.5  
548452869081677: 1, 2.5500659681053572: 1, 2.548798553113119: 1, 2.5476156268  
298618: 1, 2.5472035747293074: 1, 2.546616146301687: 1, 2.5457850615459745:  
1, 2.54394667328417: 1, 2.5432570278577975: 1, 2.542951568547263: 1, 2.542240  
800209723: 1, 2.541870956660213: 1, 2.540629850774468: 1, 2.538100060073469:  
1, 2.537493894750212: 1, 2.5320520289290323: 1, 2.5307415232728685: 1, 2.5297  
74341566591: 1, 2.526975118069478: 1, 2.5265586698293228: 1, 2.52645764195701  
53: 1, 2.526426476072916: 1, 2.525069693550488: 1, 2.524275912467874: 1, 2.52  
41339687016846: 1, 2.5223700049212963: 1, 2.5211207878096262: 1, 2.5210758399  
73921: 1, 2.5210126844054517: 1, 2.5199821030607534: 1, 2.5179567135445806:  
1, 2.5151671186862554: 1, 2.513495984353533: 1, 2.5130618680609893: 1, 2.5106  
534249066654: 1, 2.5075507565115345: 1, 2.507201208554439: 1, 2.5035577451334  
07: 1, 2.5034153628562903: 1, 2.5029022311126856: 1, 2.5028614037539376: 1,  
2.501933136472444: 1, 2.4982833211322957: 1, 2.497616876794362: 1, 2.49664119  
16272278: 1, 2.494848280806039: 1, 2.4900352917362376: 1, 2.4893830418049316:  
1, 2.485635046594663: 1, 2.485482356111894: 1, 2.4823206128754944: 1, 2.48193  
1745839503: 1, 2.4794680775350475: 1, 2.4787292806205223: 1, 2.47690352056280  
4: 1, 2.476309506961825: 1, 2.475969323424238: 1, 2.4744315434802666: 1, 2.47  
3307552590654: 1, 2.4728934013001425: 1, 2.4718778886352415: 1, 2.47134021951  
4686: 1, 2.466665776525199: 1, 2.465016005673193: 1, 2.4650027724450263: 1,  
2.463367651882188: 1, 2.4609844886666505: 1, 2.459127728647688: 1, 2.45888464  
5833307: 1, 2.4575741959560244: 1, 2.4565616363386353: 1, 2.45360543993456:  
1, 2.453163556056039: 1, 2.4526582996695616: 1, 2.452600961377953: 1, 2.45221  
9998796772: 1, 2.4516405845438687: 1, 2.451051591129413: 1, 2.44839692730360  
1: 1, 2.4480512652354087: 1, 2.447287666789577: 1, 2.4470352893102425: 1, 2.4  
45983561925909: 1, 2.4445719736279887: 1, 2.4426052443119954: 1, 2.4399307845  
964446: 1, 2.4387704098972587: 1, 2.43775655634505: 1, 2.437260780388144: 1,  
2.436919874121273: 1, 2.4356698506825554: 1, 2.435007826357198: 1, 2.43429086  
2052174: 1, 2.434288627569871: 1, 2.4330197442436985: 1, 2.4314660723908927:  
1, 2.4307287461833935: 1, 2.4279568003975753: 1, 2.4239339394143884: 1, 2.423  
2097588894494: 1, 2.4219773548330887: 1, 2.4218202005398664: 1, 2.42047604622  
8221: 1, 2.419213857583288: 1, 2.418781849596723: 1, 2.4185015467161564: 1,  
2.4162825958115786: 1, 2.416039412681658: 1, 2.4133762932820124: 1, 2.4132635  
20051742: 1, 2.4131606322957295: 1, 2.4119270744554555: 1, 2.410765591963535  
6: 1, 2.409418865822183: 1, 2.405817156723289: 1, 2.405470998051154: 1, 2.403  
6552759272074: 1, 2.400376900083287: 1, 2.4001148873231952: 1, 2.399737172525  
645: 1, 2.399510197401826: 1, 2.3968048671110878: 1, 2.396469124396718: 1, 2.  
3952031903906756: 1, 2.395108780217643: 1, 2.393634469515029: 1, 2.3929899978  
72574: 1, 2.392349724864673: 1, 2.392095596028094: 1, 2.3895172710199226: 1,  
2.387911413355165: 1, 2.387869321926015: 1, 2.3854837552138553: 1, 2.38031770  
31322704: 1, 2.3789497684971805: 1, 2.378674013277387: 1, 2.376950336225056:  
1, 2.3752386363616855: 1, 2.373126194649974: 1, 2.372916593109907: 1, 2.37242

6053679023: 1, 2.3713037436749906: 1, 2.3677761202498337: 1, 2.36658593124413  
44: 1, 2.36598749066205: 1, 2.3646140599547523: 1, 2.3633715951307313: 1, 2.3  
631960021285514: 1, 2.362654355948742: 1, 2.36247498796904: 1, 2.362206361163  
8646: 1, 2.3618664930198543: 1, 2.359597911324222: 1, 2.356838476581796: 1,  
2.356363595341795: 1, 2.3544889813019654: 1, 2.3543111029576926: 1, 2.3472068  
583679313: 1, 2.345615031017461: 1, 2.3455083682351536: 1, 2.345138771002341:  
1, 2.3445668982697243: 1, 2.342928616898407: 1, 2.3402288508798166: 1, 2.3356  
47961140897: 1, 2.333919296089879: 1, 2.3337595636306334: 1, 2.3326476550844  
6: 1, 2.3320623078222034: 1, 2.3296964982995183: 1, 2.3283438908811096: 1, 2.  
327769404584003: 1, 2.324486099957632: 1, 2.323382735330806: 1, 2.32298850599  
96542: 1, 2.3213003856131427: 1, 2.320556124688797: 1, 2.317567944004838: 1,  
2.3150107400414743: 1, 2.3146159338865746: 1, 2.314435151596901: 1, 2.3139139  
811172975: 1, 2.3102424639333656: 1, 2.310110334767819: 1, 2.309745001074183:  
1, 2.3093693093481886: 1, 2.3089536464215685: 1, 2.3060187607408564: 1, 2.305  
2736220065504: 1, 2.30480499498071: 1, 2.302839298916829: 1, 2.30149241767834  
36: 1, 2.294003590608884: 1, 2.293437755869696: 1, 2.291892140304166: 1, 2.29  
00321566488926: 1, 2.2898785363168983: 1, 2.289304168957208: 1, 2.28764635419  
2466: 1, 2.2856722584083697: 1, 2.284977374980259: 1, 2.2846284229776272: 1,  
2.283444704437413: 1, 2.281849925580031: 1, 2.280283117039155: 1, 2.272469555  
5630277: 1, 2.272062523446709: 1, 2.2687175398898347: 1, 2.2649384355732964:  
1, 2.261528770514529: 1, 2.2598439789633242: 1, 2.257943302663145: 1, 2.25702  
50598290427: 1, 2.2548919683798245: 1, 2.250836212504296: 1, 2.25082707931438  
8: 1, 2.250668851026069: 1, 2.2480156025872584: 1, 2.2460488971971393: 1, 2.2  
433326755140564: 1, 2.2409945689689397: 1, 2.2401941631791304: 1, 2.240182847  
5270946: 1, 2.239212591814302: 1, 2.2381739850944946: 1, 2.2381625192258587:  
1, 2.2368755378575518: 1, 2.235444992502204: 1, 2.233799087762679: 1, 2.22854  
86213855554: 1, 2.2281421910720045: 1, 2.2277255507591005: 1, 2.2272334895120  
287: 1, 2.2261837245367624: 1, 2.22461647435861: 1, 2.2243611638792675: 1, 2.  
222238739731972: 1, 2.21963505753031: 1, 2.219605867180967: 1, 2.216077727344  
2425: 1, 2.2120709739354183: 1, 2.210654630630541: 1, 2.210267994650048: 1,  
2.209461643510598: 1, 2.208691189963: 1, 2.208360815242066: 1, 2.208258038532  
747: 1, 2.207392755061802: 1, 2.204371923416117: 1, 2.201731195933883: 1, 2.2  
017253821023646: 1, 2.200566024826006: 1, 2.200449054094146: 1, 2.19808979249  
2036: 1, 2.196590333547758: 1, 2.192191141199017: 1, 2.189817824103931: 1, 2.  
187627270928355: 1, 2.184772089268833: 1, 2.1839492966692773: 1, 2.1816653506  
35501: 1, 2.1786951295898827: 1, 2.1772846142771694: 1, 2.1737954722987514:  
1, 2.170089395704325: 1, 2.169978441511822: 1, 2.1690372462253262: 1, 2.16665  
86750412415: 1, 2.1641412493790617: 1, 2.163611773889955: 1, 2.16232355905495  
37: 1, 2.1620618075030897: 1, 2.1522318427676903: 1, 2.150859912160984: 1, 2.  
148306100187284: 1, 2.1457172512470777: 1, 2.1404092791540976: 1, 2.140197359  
4564137: 1, 2.1396428649311483: 1, 2.1393767727430584: 1, 2.1385323735508366:  
1, 2.137433401443583: 1, 2.1366847884962263: 1, 2.1341812424631823: 1, 2.1336  
213138162425: 1, 2.1284359707481477: 1, 2.1260978339131777: 1, 2.125025729524  
54: 1, 2.1247281713458817: 1, 2.1220565914362886: 1, 2.1212574678709974: 1,  
2.1211589732431633: 1, 2.1203835725414404: 1, 2.1196741285915803: 1, 2.119233  
574944751: 1, 2.117129777244511: 1, 2.1166073506405314: 1, 2.115959632082641:  
1, 2.1089783628691494: 1, 2.1062933637418046: 1, 2.106114389502829: 1, 2.1003  
966636194766: 1, 2.0997335490567166: 1, 2.0990724933558806: 1, 2.089012249626  
9794: 1, 2.086558406235426: 1, 2.0854949441739428: 1, 2.083032380726135: 1,  
2.080822390644407: 1, 2.0792319107315893: 1, 2.0784707334556924: 1, 2.0774644  
18190011: 1, 2.072633357740014: 1, 2.0714418683417004: 1, 2.0684594159613128:  
1, 2.053529441696101: 1, 2.0457268348301447: 1, 2.0417738812291475: 1, 2.0411  
53129604513: 1, 2.0410173551142314: 1, 2.0408934711664126: 1, 2.0408803192869  
014: 1, 2.040841284268154: 1, 2.0389804740642585: 1, 2.0354905717365095: 1,  
2.0352313032609044: 1, 2.03122570328685: 1, 2.0247752581293343: 1, 2.0224188  
63402552: 1, 2.022203998132849: 1, 2.0215340911528497: 1, 2.0133647839557756:  
1, 2.0132502297925585: 1, 2.013128500301446: 1, 2.0123043418854896: 1, 2.0088

98323166728: 1, 2.0073925966659654: 1, 1.9939408925501025: 1, 1.9891994658748  
817: 1, 1.9874570796107143: 1, 1.9858313968040349: 1, 1.9838609510747374: 1,  
1.9832756406804855: 1, 1.9820186215857074: 1, 1.9817988456458813: 1, 1.97201  
85565009511: 1, 1.9712564701580726: 1, 1.970837733349584: 1, 1.96985452119162  
1: 1, 1.9623579524081578: 1, 1.9603214146229953: 1, 1.9600082710647793: 1, 1.  
9596199584915057: 1, 1.957040826359256: 1, 1.9503400753609499: 1, 1.949219317  
24851: 1, 1.9482294399994653: 1, 1.9450926728121183: 1, 1.936452841194715: 1,  
1.9310519926612: 1, 1.9299140654313762: 1, 1.9276482451878758: 1, 1.922687712  
126038: 1, 1.9218977710737408: 1, 1.9092700755453378: 1, 1.903167042117796:  
1, 1.9022386982484216: 1, 1.8988855809123546: 1, 1.8939983774394789: 1, 1.88  
92428609627265: 1, 1.8862503360627643: 1, 1.8857954153704994: 1, 1.8692370128  
152802: 1, 1.8564897178711741: 1, 1.8553226022180658: 1, 1.8225693481602918:  
1, 1.8160734064926318: 1, 1.7980075715042103: 1, 1.7956231021204738: 1, 1.79  
37441601385933: 1, 1.7861533009122217: 1, 1.7798096075942662: 1, 1.7752807162  
111215: 1, 1.7740353552439494: 1, 1.7672727984958074: 1, 1.7378972804613269:  
1, 1.689913353837599: 1})

```

In [48]: # Train a Logistic regression+Calibration model using text features which are
on-hot encoded
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])    Fit linear model with Stochastic Gradient Descent.
# predict(X)    Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_text_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_text_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)


predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)

```

```

print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
For values of alpha = 1e-05 The log loss is: 1.1766143372443885
For values of alpha = 0.0001 The log loss is: 1.1722123562885514
For values of alpha = 0.001 The log loss is: 1.4235495836242769
For values of alpha = 0.01 The log loss is: 1.7403384840951115
For values of alpha = 0.1 The log loss is: 2.0500092667630945
For values of alpha = 1 The log loss is: 2.0578269884956337

```



```

For values of best alpha = 0.0001 The train log loss is: 0.6423746250903991
For values of best alpha = 0.0001 The cross validation log loss is: 1.1722123562885514
For values of best alpha = 0.0001 The test log loss is: 1.1586735910532906

```

**Q.** Is the Text feature stable across all the data sets (Test, Train, Cross validation)?

**Ans.** Yes, it seems like!

```

In [49]: def get_intersec_text(df):
          df_text_vec = TfidfVectorizer(min_df=5, max_features=3000)
          df_text_fea = df_text_vec.fit_transform(df['TEXT'])
          df_text_features = df_text_vec.get_feature_names()

          df_text_fea_counts = df_text_fea.sum(axis=0).A1
          df_text_fea_dict = dict(zip(list(df_text_features), df_text_fea_counts))
          len1 = len(set(df_text_features))
          len2 = len(set(train_text_features) & set(df_text_features))
          return len1, len2

```

```
In [50]: len1,len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train
data")
len1,len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in
train data")
```

93.7 % of word of test data appeared in train data

92.733 % of word of Cross Validation appeared in train data

## 4. Machine Learning Models

```
In [51]: #Data preparation for ML models.

#Misc. functionns for ML models

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we willl provide the array of probabilities bel
ongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y- test_
y))/test_y.shape[0])
    plot_confusion_matrix(test_y, pred_y)
```

```
In [52]: def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)
```

```

In [53]: # this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
    text_count_vec = CountVectorizer(min_df=3)

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
                print(i, "Gene feature [{}] present in test data point [{}].f
ormat(word,yes_no))
            elif (v < fea1_len+fea2_len):
                word = var_vec.get_feature_names()[v-(fea1_len)]
                yes_no = True if word == var else False
                if yes_no:
                    word_present += 1
                    print(i, "variation feature [{}] present in test data point [
{}]].format(word,yes_no))
            else:
                word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
                yes_no = True if word in text.split() else False
                if yes_no:
                    word_present += 1
                    print(i, "Text feature [{}] present in test data point [{}].f
ormat(word,yes_no))

    print("Out of the top ",no_features," features ", word_present, "are prese
nt in query point")

```

## Stacking the three types of features

```

In [54]: # merging gene, variance and text features

# building train, test and cross validation data sets
# a = [[1, 2],
#      [3, 4]]
# b = [[4, 5],
#      [6, 7]]
# hstack(a, b) = [[1, 2, 4, 5],
#                [ 3, 4, 6, 7]]

train_gene_var_onehotCoding = hstack((train_gene_feature_onehotCoding, train_variation_feature_onehotCoding))
test_gene_var_onehotCoding = hstack((test_gene_feature_onehotCoding, test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding, cv_variation_feature_onehotCoding))

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotCoding)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCoding)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).tocsr()
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding = np.hstack((train_gene_feature_responseCoding, train_variation_feature_responseCoding))
test_gene_var_responseCoding = np.hstack((test_gene_feature_responseCoding, test_variation_feature_responseCoding))
cv_gene_var_responseCoding = np.hstack((cv_gene_feature_responseCoding, cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding, train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding))
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))

```



```
In [55]: print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_
_onehotCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_o
nehotCoding.shape)
print("(number of data points * number of features) in cross validation data
=", cv_x_onehotCoding.shape)
```

```
One hot encoding features :
(number of data points * number of features) in train data = (2124, 5197)
(number of data points * number of features) in test data = (665, 5197)
(number of data points * number of features) in cross validation data = (532,
5197)
```

```
In [56]: print(" Response encoding features :")
print("(number of data points * number of features) in train data = ", train_x
_responseCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_r
esponseCoding.shape)
print("(number of data points * number of features) in cross validation data
=", cv_x_responseCoding.shape)
```

```
Response encoding features :
(number of data points * number of features) in train data = (2124, 27)
(number of data points * number of features) in test data = (665, 27)
(number of data points * number of features) in cross validation data = (532,
27)
```

## Feature Engineering on one-hot encoded features

```
In [57]: train_x_onehotCodingFE=np.sqrt(train_x_onehotCoding)
test_x_onehotCodingFE=np.sqrt(test_x_onehotCoding)
cv_x_onehotCodingFE=np.sqrt(cv_x_onehotCoding)
```

### 4.1. Base Line Model

#### 4.1.1. Naive Bayes

##### 4.1.1.1. Hyper parameter tuning

```

In [58]: # find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to X, y
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = MultinomialNB(alpha=i)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(np.log10(alpha), cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (np.log10(alpha[i]), cv_log_error_array[i]))

```

```
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for alpha = 1e-05
Log Loss : 1.2442427651450625
for alpha = 0.0001
Log Loss : 1.2454826518681834
for alpha = 0.001
Log Loss : 1.2423007584053438
for alpha = 0.1
Log Loss : 1.2656989120604816
for alpha = 1
Log Loss : 1.3155405657141117
for alpha = 10
Log Loss : 1.4817120025323702
for alpha = 100
Log Loss : 1.4509086811280916
for alpha = 1000
Log Loss : 1.4194057044076456

```



For values of best alpha = 0.001 The train log loss is: 0.5827010882734858  
 For values of best alpha = 0.001 The cross validation log loss is: 1.2423007584053438  
 For values of best alpha = 0.001 The test log loss is: 1.2647146781626806

#### 4.1.1.2. Testing the model with best hyper paramters

```

In [59]: # find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight])    Fit Naive Bayes classifier according to X, y
# predict(X)    Perform classification on an array of test vectors X.
# predict_log_proba(X)    Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
# -----

clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
# to avoid rounding error while multiplying probabilities we use log-probability estimates
print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("Number of misclassified point :", np.count_nonzero((sig_clf.predict(cv_x_onehotCoding) - cv_y))/cv_y.shape[0])
plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding.toarray()))

```


Log Loss : 1.2423007584053438

Number of missclassified point : 0.40037593984962405


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



#### 4.1.1.3. Feature Importance, Correctly classified point

```
In [60]: test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices=np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0811 0.0627 0.0133 0.0758 0.0384 0.0401 0.  
679 0.0055 0.0042]]

Actual Class : 7

-----

Out of the top 100 features 0 are present in query point

#### 4.1.1.4. Feature Importance, Incorrectly classified point

```
In [61]: test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 2

Predicted Class Probabilities: [[0.0718 0.6318 0.0143 0.0814 0.0412 0.0434 0.  
1057 0.0059 0.0045]]

Actual Class : 2

-----

Out of the top 100 features 0 are present in query point

## 4.2. K Nearest Neighbour Classification

### 4.2.1. Hyper parameter tuning

```

In [62]: # find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----

# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-neighbors-geometric-intuition-with-a-toy-example-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])           Get parameters for this estimator.
# predict(X)                    Predict the target of new samples.
# predict_proba(X)              Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")

```



```
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for alpha = 5
Log Loss : 1.0969303646472561
for alpha = 11
Log Loss : 1.0851016136985003
for alpha = 15
Log Loss : 1.0970109480905619
for alpha = 21
Log Loss : 1.0967950051940067
for alpha = 31
Log Loss : 1.0897571502715409
for alpha = 41
Log Loss : 1.0816610736908696
for alpha = 51
Log Loss : 1.081632841418107
for alpha = 99
Log Loss : 1.1126261578804915

```



For values of best alpha = 51 The train log loss is: 0.8691792565078065  
 For values of best alpha = 51 The cross validation log loss is: 1.081632841418107  
 For values of best alpha = 51 The test log loss is: 1.0883792266791434

#### 4.2.2. Testing the model with best hyper paramters


```
In [63]: # find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-neighbors-geometric-intuition-with-a-toy-example-1/
#-----
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y, cv_x_responseCoding, cv_y, clf)
```

Log loss : 1.081632841418107

Number of mis-classified points : 0.37030075187969924

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 4.2.3.Sample Query point -1

```
In [64]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
predicted_cls = sig_clf.predict(test_x_responseCoding[0].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1,
-1), alpha[best_alpha])
print("The ",alpha[best_alpha]," nearest neighbours of the test points belongs
to classes",train_y[neighbors[1][0]])
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]]))

Predicted Class : 1
Actual Class : 7
The 51 nearest neighbours of the test points belongs to classes [7 6 1 6 1
1 1 6 1 7 7 7 7 7 2 7 2 7 7 2 7 7 1 7 2 6 2 2 7 7 6 7 2 7 2 7 7
2 2 7 7 4 7 7 7 7 7 7 7 1]
Fequency of nearest points : Counter({7: 28, 2: 10, 1: 7, 6: 5, 4: 1})
```

#### 4.2.4. Sample Query Point-2

```
In [65]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 100

predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshap
e(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1,
-1), alpha[best_alpha])
print("the k value for knn is",alpha[best_alpha],"and the nearest neighbours o
f the test points belongs to classes",train_y[neighbors[1][0]])
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]]))

Predicted Class : 2
Actual Class : 2
the k value for knn is 51 and the nearest neighbours of the test points belon
gs to classes [2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 7 2 2 2 2 2 2 2 2 2 2 2
2 7 7 7 2 7
2 2 2 2 2 7 7 2 2 7 7 2 7]
Fequency of nearest points : Counter({2: 41, 7: 10})
```

### 4.3. Logistic Regression

## With Count Vectorizer and Unigram and bigram

```
In [66]: # one-hot encoding of Gene feature.
gene_vectorizer_lr = CountVectorizer()
train_gene_feature_onehotCoding_lr = gene_vectorizer_lr.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding_lr = gene_vectorizer_lr.transform(test_df['Gene'])
cv_gene_feature_onehotCoding_lr = gene_vectorizer_lr.transform(cv_df['Gene'])

# one-hot encoding of variation feature.
variation_vectorizer_lr = CountVectorizer()
train_variation_feature_onehotCoding_lr = variation_vectorizer_lr.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding_lr = variation_vectorizer_lr.transform(test_df['Variation'])
cv_variation_feature_onehotCoding_lr = variation_vectorizer_lr.transform(cv_df['Variation'])

text_vectorizer_lr = CountVectorizer(ngram_range=(1, 4))
train_text_feature_onehotCoding_lr = text_vectorizer_lr.fit_transform(train_df['TEXT'])
# don't forget to normalize every feature
train_text_feature_onehotCoding_lr = normalize(train_text_feature_onehotCoding_lr, axis=0)

test_text_feature_onehotCoding_lr = text_vectorizer_lr.transform(test_df['TEXT'])
test_text_feature_onehotCoding_lr = normalize(test_text_feature_onehotCoding_lr, axis=0)

cv_text_feature_onehotCoding_lr = text_vectorizer_lr.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding_lr = normalize(cv_text_feature_onehotCoding_lr, axis=0)
```

**stacking all the features(gene, variations, text of one-hot encoded)**

```
In [67]: train_gene_var_onehotCoding_lr = hstack((train_gene_feature_onehotCoding_lr, train_variation_feature_onehotCoding_lr))
test_gene_var_onehotCoding_lr = hstack((test_gene_feature_onehotCoding_lr, test_variation_feature_onehotCoding_lr))
cv_gene_var_onehotCoding_lr = hstack((cv_gene_feature_onehotCoding_lr, cv_variation_feature_onehotCoding_lr))

train_x_onehotCoding_lr = hstack((train_gene_var_onehotCoding_lr, train_text_feature_onehotCoding_lr)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding_lr = hstack((test_gene_var_onehotCoding_lr, test_text_feature_onehotCoding_lr)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding_lr = hstack((cv_gene_var_onehotCoding_lr, cv_text_feature_onehotCoding_lr)).tocsr()
cv_y = np.array(list(cv_df['Class']))
```

```
In [68]: print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_onehotCoding_lr.shape)
print("(number of data points * number of features) in test data = ", test_x_onehotCoding_lr.shape)
print("(number of data points * number of features) in cross validation data = ", cv_x_onehotCoding_lr.shape)
```

```
One hot encoding features :
(number of data points * number of features) in train data = (2124, 11915609)
(number of data points * number of features) in test data = (665, 11915609)
(number of data points * number of features) in cross validation data = (532, 11915609)
```

### 4.3.1. With Class balancing

#### 4.3.1.1. Hyper paramter tuning

```

In [69]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])    Fit linear model with Stochastic Gradient Descent.
# predict(X)    Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding_lr, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding_lr, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding_lr)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):

```



```
ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty=
'12', loss='log', random_state=42)
clf.fit(train_x_onehotCoding_lr, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding_lr, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss i
s:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss i
s:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for alpha = 1e-06
Log Loss : 1.8201104916799518
for alpha = 1e-05
Log Loss : 1.6203771586952185
for alpha = 0.0001
Log Loss : 1.5875493680791593
for alpha = 0.001
Log Loss : 1.403376474810247
for alpha = 0.01
Log Loss : 1.2423272309940365
for alpha = 0.1
Log Loss : 1.2401531939506085
for alpha = 1
Log Loss : 1.2532748766119728
for alpha = 10
Log Loss : 1.244574367442384
for alpha = 100
Log Loss : 1.2491176828345538

```



For values of best alpha = 0.1 The train log loss is: 0.5540673088073816  
 For values of best alpha = 0.1 The cross validation log loss is: 1.2401531939506085  
 For values of best alpha = 0.1 The test log loss is: 1.2425756049800392

#### 4.3.1.2. Testing the model with best hyper paramters

```

In [70]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])    Fit linear model with Stochastic Gradient Descent.
# predict(X)    Predict class labels for samples in X.


#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding_lr, train_y, cv_x_onehotCoding_lr, cv_y, clf)

```


Log loss : 1.2401531939506085

Number of mis-classified points : 0.3890977443609023


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 4.3.1.3. Feature Importance

```

In [71]: #def get_imp_feature_names(text, indices, removed_ind = []):
#         word_present = 0
#         tabulte_list = []
#         incresingorder_ind = 0
#         for i in indices:
#             if i < train_gene_feature_onehotCoding.shape[1]:
#                 tabulte_list.append([incresingorder_ind, "Gene", "Yes"])
#             elif i < 18:
#                 tabulte_list.append([incresingorder_ind, "Variation", "Yes"])
#             if ((i > 17) & (i not in removed_ind)) :
#                 word = train_text_features[i]
#                 yes_no = True if word in text.split() else False
#                 if yes_no:
#                     word_present += 1
#                 tabulte_list.append([incresingorder_ind, train_text_features[i], y
# es_no])
#                 incresingorder_ind += 1
#         print(word_present, "most important features are present in our query poi
# nt")
#         print("-"*50)
#         print("The features that are most important of the ", predicted_cls[0], " c
# lass:")
#         print(tabulate(tabulte_list, headers=["Index", 'Feature name', 'Present o
# r Not']))
#         # this function will be used just for naive bayes
#         # for the given indices, we will print the name of the features
#         # and we will check whether the feature present in the test point text or not

def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = TfidfVectorizer()
    var_count_vec = TfidfVectorizer()
    text_count_vec = TfidfVectorizer(ngram_range=(1, 4))

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_count_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
            print(i, "Gene feature [{}] present in test data point [{}]"
. format(word,yes_no))
        elif (v < fea1_len+fea2_len):
            word = var_vec.get_feature_names()[v-(fea1_len)]
            yes_no = True if word == var else False
            if yes_no:
                word_present += 1
            print(i, "variation feature [{}] present in test data point [
{}]"
. format(word,yes_no))

```

```

else:
    word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
    yes_no = True if word in text.split() else False
    if yes_no:
        word_present += 1
        print(i, "Text feature [{}] present in test data point [{}].".f
ormat(word,yes_no))

    print("Out of the top ",no_features," features ", word_present, "are prese
nt in query point")

```

#### 4.3.1.3.1. Correctly Classified point

```

In [72]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty=
'12', loss='log', random_state=42)
clf.fit(train_x_onehotCoding_lr,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding_lr[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding_lr[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
#indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,:no_feature]
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)

Predicted Class : 7
Predicted Class Probabilities: [[0.1836 0.1433 0.0285 0.1797 0.077 0.0788 0.
2936 0.005 0.0103]]
Actual Class : 7
-----
Out of the top 500 features 0 are present in query point

```

#### 4.3.1.3.2. Incorrectly Classified point

```
In [73]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding_lr[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding_lr[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 2

Predicted Class Probabilities: [[0.0122 0.9412 0.0022 0.0096 0.006 0.003 0.0185 0.005 0.0024]]

Actual Class : 2

-----

32 Text feature [wild] present in test data point [True]  
33 Text feature [type] present in test data point [True]  
36 Text feature [predicted] present in test data point [True]  
38 Text feature [results] present in test data point [True]  
41 Text feature [whereas] present in test data point [True]  
43 Text feature [acid] present in test data point [True]  
44 Text feature [used] present in test data point [True]  
46 Text feature [amino] present in test data point [True]  
47 Text feature [likely] present in test data point [True]  
49 Text feature [determine] present in test data point [True]  
50 Text feature [sequence] present in test data point [True]  
54 Text feature [expected] present in test data point [True]  
61 Text feature [data] present in test data point [True]  
62 Text feature [thus] present in test data point [True]  
63 Text feature [whether] present in test data point [True]  
64 Text feature [although] present in test data point [True]  
65 Text feature [shown] present in test data point [True]  
66 Text feature [showed] present in test data point [True]  
67 Text feature [previously] present in test data point [True]  
68 Text feature [introduction] present in test data point [True]  
70 Text feature [indicate] present in test data point [True]  
72 Text feature [therefore] present in test data point [True]  
73 Text feature [independent] present in test data point [True]  
76 Text feature [containing] present in test data point [True]  
77 Text feature [additional] present in test data point [True]  
82 Text feature [also] present in test data point [True]  
84 Text feature [mutant] present in test data point [True]  
85 Text feature [using] present in test data point [True]  
86 Text feature [addition] present in test data point [True]  
87 Text feature [either] present in test data point [True]  
88 Text feature [provide] present in test data point [True]  
90 Text feature [well] present in test data point [True]  
93 Text feature [indicated] present in test data point [True]  
95 Text feature [result] present in test data point [True]  
102 Text feature [found] present in test data point [True]  
103 Text feature [changes] present in test data point [True]  
105 Text feature [analysis] present in test data point [True]  
112 Text feature [full] present in test data point [True]  
113 Text feature [important] present in test data point [True]  
115 Text feature [analyzed] present in test data point [True]  
116 Text feature [critical] present in test data point [True]  
117 Text feature [mutations] present in test data point [True]  
118 Text feature [proteins] present in test data point [True]  
119 Text feature [one] present in test data point [True]  
121 Text feature [50] present in test data point [True]  
124 Text feature [may] present in test data point [True]  
125 Text feature [associated] present in test data point [True]  
127 Text feature [domain] present in test data point [True]  
130 Text feature [possible] present in test data point [True]  
134 Text feature [resulting] present in test data point [True]  
137 Text feature [assays] present in test data point [True]  
138 Text feature [methods] present in test data point [True]



139 Text feature [directly] present in test data point [True]  
141 Text feature [reduced] present in test data point [True]  
146 Text feature [change] present in test data point [True]  
147 Text feature [lower] present in test data point [True]  
148 Text feature [figure] present in test data point [True]  
151 Text feature [based] present in test data point [True]  
152 Text feature [form] present in test data point [True]  
153 Text feature [expressed] present in test data point [True]  
156 Text feature [fact] present in test data point [True]  
159 Text feature [cells] present in test data point [True]  
160 Text feature [known] present in test data point [True]  
161 Text feature [substitutions] present in test data point [True]  
163 Text feature [note] present in test data point [True]  
164 Text feature [highly] present in test data point [True]  
166 Text feature [vitro] present in test data point [True]  
168 Text feature [however] present in test data point [True]  
170 Text feature [binding] present in test data point [True]  
179 Text feature [levels] present in test data point [True]  
181 Text feature [several] present in test data point [True]  
184 Text feature [identified] present in test data point [True]  
189 Text feature [performed] present in test data point [True]  
193 Text feature [effects] present in test data point [True]  
194 Text feature [mutation] present in test data point [True]  
195 Text feature [represent] present in test data point [True]  
196 Text feature [general] present in test data point [True]  
197 Text feature [10] present in test data point [True]  
198 Text feature [table] present in test data point [True]  
199 Text feature [cancer] present in test data point [True]  
201 Text feature [contribute] present in test data point [True]  
204 Text feature [reported] present in test data point [True]  
205 Text feature [members] present in test data point [True]  
209 Text feature [indicates] present in test data point [True]  
210 Text feature [residues] present in test data point [True]  
213 Text feature [approximately] present in test data point [True]  
214 Text feature [least] present in test data point [True]  
217 Text feature [suggested] present in test data point [True]  
221 Text feature [could] present in test data point [True]  
227 Text feature [presence] present in test data point [True]  
230 Text feature [30] present in test data point [True]  
238 Text feature [would] present in test data point [True]  
239 Text feature [different] present in test data point [True]  
241 Text feature [genetic] present in test data point [True]  
242 Text feature [cause] present in test data point [True]  
245 Text feature [respectively] present in test data point [True]  
247 Text feature [large] present in test data point [True]  
248 Text feature [within] present in test data point [True]  
249 Text feature [absence] present in test data point [True]  
254 Text feature [derived] present in test data point [True]  
255 Text feature [comparison] present in test data point [True]  
259 Text feature [25] present in test data point [True]  
265 Text feature [selected] present in test data point [True]  
267 Text feature [predict] present in test data point [True]  
270 Text feature [obtained] present in test data point [True]  
274 Text feature [taken] present in test data point [True]  
277 Text feature [26] present in test data point [True]  
290 Text feature [available] present in test data point [True]  
293 Text feature [possibility] present in test data point [True]

296 Text feature [confirmed] present in test data point [True]  
301 Text feature [fold] present in test data point [True]  
303 Text feature [studies] present in test data point [True]  
304 Text feature [even] present in test data point [True]  
305 Text feature [except] present in test data point [True]  
306 Text feature [24] present in test data point [True]  
310 Text feature [position] present in test data point [True]  
311 Text feature [allows] present in test data point [True]  
315 Text feature [present] present in test data point [True]  
320 Text feature [identify] present in test data point [True]  
321 Text feature [support] present in test data point [True]  
322 Text feature [given] present in test data point [True]  
324 Text feature [observation] present in test data point [True]  
331 Text feature [approach] present in test data point [True]  
332 Text feature [forms] present in test data point [True]  
333 Text feature [cell] present in test data point [True]  
336 Text feature [strong] present in test data point [True]  
341 Text feature [yet] present in test data point [True]  
342 Text feature [according] present in test data point [True]  
345 Text feature [caused] present in test data point [True]  
348 Text feature [relevant] present in test data point [True]  
349 Text feature [specific] present in test data point [True]  
350 Text feature [include] present in test data point [True]  
352 Text feature [useful] present in test data point [True]  
359 Text feature [frequently] present in test data point [True]  
362 Text feature [significantly] present in test data point [True]  
365 Text feature [similarly] present in test data point [True]  
375 Text feature [recently] present in test data point [True]  
378 Text feature [remaining] present in test data point [True]  
384 Text feature [35] present in test data point [True]  
385 Text feature [allow] present in test data point [True]  
387 Text feature [low] present in test data point [True]  
389 Text feature [grown] present in test data point [True]  
391 Text feature [many] present in test data point [True]  
397 Text feature [single] present in test data point [True]  
400 Text feature [standard] present in test data point [True]  
403 Text feature [model] present in test data point [True]  
404 Text feature [assessment] present in test data point [True]  
409 Text feature [little] present in test data point [True]  
415 Text feature [proposed] present in test data point [True]  
418 Text feature [resulted] present in test data point [True]  
420 Text feature [including] present in test data point [True]  
424 Text feature [detected] present in test data point [True]  
425 Text feature [assessed] present in test data point [True]  
435 Text feature [involved] present in test data point [True]  
436 Text feature [less] present in test data point [True]  
437 Text feature [basis] present in test data point [True]  
441 Text feature [exception] present in test data point [True]  
452 Text feature [structural] present in test data point [True]  
453 Text feature [high] present in test data point [True]  
458 Text feature [panel] present in test data point [True]  
459 Text feature [impair] present in test data point [True]  
463 Text feature [40] present in test data point [True]  
466 Text feature [combined] present in test data point [True]  
476 Text feature [difficult] present in test data point [True]  
478 Text feature [31] present in test data point [True]  
483 Text feature [published] present in test data point [True]

484 Text feature [comprehensive] present in test data point [True]  
487 Text feature [abrogate] present in test data point [True]  
489 Text feature [reduction] present in test data point [True]  
496 Text feature [relatively] present in test data point [True]  
497 Text feature [29] present in test data point [True]  
499 Text feature [might] present in test data point [True]  
Out of the top 500 features 172 are present in query point

## 4.3.2. Without Class balancing

### 4.3.2.1. Hyper paramter tuning

```

In [74]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])      Fit the calibrated model
# get_params([deep])      Get parameters for this estimator.
# predict(X)      Predict the target of new samples.
# predict_proba(X)      Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding_lr, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding_lr, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding_lr)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()

```

```
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_
state=42)
clf.fit(train_x_onehotCoding_lr, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding_lr, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss i
s:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding_lr)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss i
s:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for alpha = 1e-06
Log Loss : 1.774894857014709
for alpha = 1e-05
Log Loss : 1.5578290286106147
for alpha = 0.0001
Log Loss : 1.5544148761287524
for alpha = 0.001
Log Loss : 1.4115997798122784
for alpha = 0.01
Log Loss : 1.2223735845597772
for alpha = 0.1
Log Loss : 1.2294773325874306
for alpha = 1
Log Loss : 1.2379790995332929

```



For values of best alpha = 0.01 The train log loss is: 0.6068200320716836  
 For values of best alpha = 0.01 The cross validation log loss is: 1.2223735845597772  
 For values of best alpha = 0.01 The test log loss is: 1.235224292000013

#### 4.3.2.2. Testing model with best hyper parameters

```
In [75]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])    Fit linear model with Stochastic Gradient Descent.
# predict(X)    Predict class labels for samples in X.


#-----
# video link:
#-----

clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding_lr, train_y, cv_x_onehotCoding_lr, cv_y, clf)
```


Log loss : 1.2223735845597772

Number of mis-classified points : 0.40977443609022557


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



#### 4.3.2.3. Feature Importance, Correctly Classified point



```
In [76]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_
state=42)
clf.fit(train_x_onehotCoding_lr,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding_lr[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding_lr[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.1797 0.1384 0.0365 0.1727 0.0877 0.0913 0.  
2744 0.0055 0.0138]]

Actual Class : 7

-----

Out of the top 500 features 0 are present in query point

#### 4.3.2.4. Feature Importance, Inorrectly Classified point

```
In [77]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding_lr[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding_lr[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 2

Predicted Class Probabilities: [[0.0179 0.8893 0.0069 0.0126 0.0096 0.0048 0.  
0507 0.006 0.0022]]

Actual Class : 2

-----

Out of the top 500 features 0 are present in query point

## 4.4. Linear Support Vector Machines

### 4.4.1. Hyper paramter tuning

```

In [78]: # read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True,
# probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given training data.
# predict(X)    Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default parameters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
    #     clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))

```

```
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for C = 1e-05
Log Loss : 1.0084816711655122
for C = 0.0001
Log Loss : 1.0211186485835775
for C = 0.001
Log Loss : 1.0424249526918998
for C = 0.01
Log Loss : 1.2923698826699872
for C = 0.1
Log Loss : 1.8817742008573548
for C = 1
Log Loss : 1.9360273420627463
for C = 10
Log Loss : 1.9360251971882068
for C = 100
Log Loss : 1.9360242795186076

```



For values of best alpha = 1e-05 The train log loss is: 0.41723500769698235  
 For values of best alpha = 1e-05 The cross validation log loss is: 1.0084816711655122  
 For values of best alpha = 1e-05 The test log loss is: 1.0867171998905403

#### 4.4.2. Testing model with best hyper parameters

```
In [79]: # read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True,
# probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)


# Some of methods of SVM()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given training data.
# predict(X)    Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# clf = SVC(C=alpha[best_alpha],kernel='linear',probability=True, class_weight='balanced')
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42,class_weight='balanced')
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_onehotCoding,cv_y, clf)
```


Log loss : 1.0084816711655122

Number of mis-classified points : 0.32142857142857145


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 4.3.3. Feature Importance

#### 4.3.3.1. For Correctly classified point

```
In [80]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
# test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation'].iloc[test_point_index], no_feature)
```

```
Predicted Class : 7
Predicted Class Probabilities: [[0.1045 0.0585 0.0056 0.0072 0.0537 0.0231 0.7386 0.0079 0.0009]]
Actual Class : 7
-----
Out of the top 500 features 0 are present in query point
```

#### 4.3.3.2. For Incorrectly classified point

```
In [81]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation'].iloc[test_point_index], no_feature)
```

```
Predicted Class : 2
Predicted Class Probabilities: [[6.760e-02 8.167e-01 1.300e-03 3.200e-03 5.610e-02 1.910e-02 2.900e-02 6.700e-03 3.000e-04]]
Actual Class : 2
-----
Out of the top 500 features 0 are present in query point
```

## 4.5 Random Forest Classifier

### **4.5.1. Hyper paramter tuning (With One hot Encoding)**



```

In [82]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', m
ax_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_l
eaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_s
tate=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given train
ing data.
# predict(X)    Perform classification on samples in X.
# predict_proba (X)    Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigm
oid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators = ", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_dep
th=j, random_state=42, n_jobs=-1)
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.cla
sses_, eps=1e-15))

```

```

print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)), (features[i],c
v_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion=
'gini', max_depth=max_depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train
log loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross
validation log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1
e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test l
og loss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```
for n_estimators = 100 and max depth = 5
Log Loss : 1.2281311521625655
for n_estimators = 100 and max depth = 10
Log Loss : 1.2016759360832534
for n_estimators = 200 and max depth = 5
Log Loss : 1.2175515228590474
for n_estimators = 200 and max depth = 10
Log Loss : 1.1925979376924802
for n_estimators = 500 and max depth = 5
Log Loss : 1.219568165490032
for n_estimators = 500 and max depth = 10
Log Loss : 1.1891083251834829
for n_estimators = 1000 and max depth = 5
Log Loss : 1.215188954131631
for n_estimators = 1000 and max depth = 10
Log Loss : 1.1871364831297526
for n_estimators = 2000 and max depth = 5
Log Loss : 1.2146143588782397
for n_estimators = 2000 and max depth = 10
Log Loss : 1.1853708880800045
For values of best estimator = 2000 The train log loss is: 0.529806763228517
7
For values of best estimator = 2000 The cross validation log loss is: 1.1853
708880800042
For values of best estimator = 2000 The test log loss is: 1.18208984475796
```

#### 4.5.2. Testing model with best hyper parameters (One Hot Encoding)

```

In [83]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', m
ax_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_l
eaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_s
tate=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given train
ing data.
# predict(X)    Perform classification on samples in X.
# predict_proba (X)    Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion=
'gini', max_depth=max_depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_onehotCod
ing,cv_y, clf)

```


Log loss : 1.1853708880800045

Number of mis-classified points : 0.3890977443609023


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 4.5.3. Feature Importance

#### 4.5.3.1. Correctly Classified point

```
In [84]: # test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion=
'gini', max_depth=max_depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_ind
ex],test_df['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_poin
t_index], no_feature)

Predicted Class : 7
Predicted Class Probabilities: [[0.1483 0.2316 0.0227 0.1459 0.06 0.0604 0.
3133 0.0084 0.0094]]
Actual Class : 7
-----
Out of the top 100 features 0 are present in query point
```

#### 4.5.3.2. Inorrectly Classified point

```
In [85]: test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_ind
ex],test_df['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_poin
t_index], no_feature)

Predicted Class : 2
Predicted Class Probabilities: [[0.0374 0.8232 0.0112 0.0237 0.0305 0.0276 0.
0379 0.0038 0.0047]]
Actual Class : 2
-----
Out of the top 100 features 0 are present in query point
```

### **4.5.3. Hyper paramter tuning (With Response Coding)**

```

In [86]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', m
ax_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_l
eaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_s
tate=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given train
ing data.
# predict(X)    Perform classification on samples in X.
# predict_proba (X)    Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/
lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stabl
e/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigm
oid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])    Fit the calibrated model
# get_params([deep])    Get parameters for this estimator.
# predict(X)    Predict the target of new samples.
# predict_proba(X)    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators = ", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_dep
th=j, random_state=42, n_jobs=-1)
        clf.fit(train_x_responseCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_responseCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.cla
sses_, eps=1e-15))

```



```

        print("Log Loss :",log_loss(cv_y, sig_clf_probs))
    '''
fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/4)],max_depth[int(i%4)],str(txt)), (features[i],c
v_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion=
'gini', max_depth=max_depth[int(best_alpha%4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The train log
loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cross vali
dation log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15
))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The test log l
oss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 10 and max depth = 2
Log Loss : 2.061724756653649
for n_estimators = 10 and max depth = 3
Log Loss : 1.7910329172541843
for n_estimators = 10 and max depth = 5
Log Loss : 1.4895860644462515
for n_estimators = 10 and max depth = 10
Log Loss : 1.6284775889672949
for n_estimators = 50 and max depth = 2
Log Loss : 1.6817896209283998
for n_estimators = 50 and max depth = 3
Log Loss : 1.4416309015498987
for n_estimators = 50 and max depth = 5
Log Loss : 1.3743188677368006
for n_estimators = 50 and max depth = 10
Log Loss : 1.6418261536253445
for n_estimators = 100 and max depth = 2
Log Loss : 1.5642606651228206
for n_estimators = 100 and max depth = 3
Log Loss : 1.4474642777930014
for n_estimators = 100 and max depth = 5
Log Loss : 1.296283409073802
for n_estimators = 100 and max depth = 10
Log Loss : 1.707053409615415
for n_estimators = 200 and max depth = 2
Log Loss : 1.584006280908212
for n_estimators = 200 and max depth = 3
Log Loss : 1.4596795181637647
for n_estimators = 200 and max depth = 5
Log Loss : 1.3515525796758852
for n_estimators = 200 and max depth = 10
Log Loss : 1.67543766818155
for n_estimators = 500 and max depth = 2
Log Loss : 1.6037968514884746
for n_estimators = 500 and max depth = 3
Log Loss : 1.491675068311772
for n_estimators = 500 and max depth = 5
Log Loss : 1.3527340554950347
for n_estimators = 500 and max depth = 10
Log Loss : 1.6866252632477592
for n_estimators = 1000 and max depth = 2
Log Loss : 1.588058323618764
for n_estimators = 1000 and max depth = 3
Log Loss : 1.5009066626880303
for n_estimators = 1000 and max depth = 5
Log Loss : 1.354320213853695
for n_estimators = 1000 and max depth = 10
Log Loss : 1.6762314802392537
For values of best alpha = 100 The train log loss is: 0.0690817567521354
For values of best alpha = 100 The cross validation log loss is: 1.296283409
0738018
For values of best alpha = 100 The test log loss is: 1.2929752726522183

```

#### 4.5.4. Testing model with best hyper parameters (Response Coding)

```

In [87]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', m
ax_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_l
eaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_s
tate=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])    Fit the SVM model according to the given train
ing data.
# predict(X)    Perform classification on samples in X.
# predict_proba (X)    Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----


clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)], n_estimat
ors=alpha[int(best_alpha/4)], criterion='gini', max_features='auto', random_sta
te=42)
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,cv_x_respons
eCoding,cv_y, clf)

```

Log loss : 1.296283409073802

Number of mis-classified points : 0.4191729323308271


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## 4.5.5. Feature Importance

#### **4.5.5.1. Correctly Classified point**

```
In [88]: clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion=
'gini', max_depth=max_depth[int(best_alpha%4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
no_feature = 27
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")
```

Predicted Class : 6

Predicted Class Probabilities: [[0.0682 0.167 0.1092 0.062 0.096 0.1889 0.1747 0.0791 0.0547]]

Actual Class : 7

-----  
Variation is important feature  
Variation is important feature  
Variation is important feature  
Variation is important feature  
Gene is important feature  
Variation is important feature  
Text is important feature  
Variation is important feature  
Text is important feature  
Gene is important feature  
Text is important feature  
Text is important feature  
Text is important feature  
Text is important feature  
Gene is important feature  
Variation is important feature  
Gene is important feature  
Gene is important feature  
Text is important feature  
Variation is important feature  
Gene is important feature  
Variation is important feature  
Text is important feature  
Text is important feature  
Text is important feature  
Gene is important feature  
Gene is important feature  
Gene is important feature

#### 4.5.5.2. Incorrectly Classified point

```
In [89]: test_point_index = 100
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")
```

Predicted Class : 2

Predicted Class Probabilities: [[0.0251 0.7726 0.0309 0.0272 0.017 0.035 0.0373 0.0286 0.0263]]

Actual Class : 2

```
-----
Variation is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Text is important feature
Variation is important feature
Text is important feature
Gene is important feature
Text is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Variation is important feature
Gene is important feature
Gene is important feature
Text is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Gene is important feature
```

## 4.7 Stack the models



### 4.7.1 testing with hyper parameter tuning

```

In [90]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight])      Fit the SVM model according to the given training data.
# predict(X)      Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])      Fit the SVM model according to the given training data.

```

```

# predict(X)      Perform classification on samples in X.
# predict_proba (X)      Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.applidaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight='balanced', random_state=0)
clf1.fit(train_x_onehotCoding, train_y)
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")

clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight='balanced', random_state=0)
clf2.fit(train_x_onehotCoding, train_y)
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")

clf3 = MultinomialNB(alpha=0.001)
clf3.fit(train_x_onehotCoding, train_y)
sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")

sig_clf1.fit(train_x_onehotCoding, train_y)
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding))))
sig_clf2.fit(train_x_onehotCoding, train_y)
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding))))
sig_clf3.fit(train_x_onehotCoding, train_y)
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding))))
print("-"*50)
alpha = [0.0001,0.001,0.01,0.1,1,10]
best_alpha = 999
for i in alpha:
    lr = LogisticRegression(C=i)
    sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_probas=True)
    sclf.fit(train_x_onehotCoding, train_y)
    print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" % (i, log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))))
    log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
    if best_alpha > log_error:
        best_alpha = log_error

```

Logistic Regression : Log Loss: 1.09  
Support vector machines : Log Loss: 1.94  
Naive Bayes : Log Loss: 1.24

-----  
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 1.817  
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 1.715  
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.320  
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.243  
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.598  
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.977

#### 4.7.2 testing the model with the best hyper parameters

```
In [91]: lr = LogisticRegression(C=0.1)
sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_probabilities=True)
sclf.fit(train_x_onehotCoding, train_y)

log_error = log_loss(train_y, sclf.predict_proba(train_x_onehotCoding))
print("Log loss (train) on the stacking classifier :", log_error)

log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
print("Log loss (CV) on the stacking classifier :", log_error)

log_error = log_loss(test_y, sclf.predict_proba(test_x_onehotCoding))
print("Log loss (test) on the stacking classifier :", log_error)

print("Number of misclassified points :", np.count_nonzero((sclf.predict(test_x_onehotCoding) - test_y).astype(int)))
plot_confusion_matrix(test_y=test_y, predict_y=sclf.predict(test_x_onehotCoding))
```


Log loss (train) on the stacking classifier : 0.4597113477737592

Log loss (CV) on the stacking classifier : 1.2434078280255831


Log loss (test) on the stacking classifier : 1.2759942326172031

Number of missclassified point : 0.3849624060150376

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



### 4.7.3 Maximum Voting classifier

```
In [92]: #Refer:http://scikit-Learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier.html  
from sklearn.ensemble import VotingClassifier  
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf',  
    sig_clf3)], voting='soft')  
vclf.fit(train_x_onehotCoding, train_y)  
print("Log loss (train) on the VotingClassifier :", log_loss(train_y, vclf.predict_proba(train_x_onehotCoding)))  
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y, vclf.predict_proba(cv_x_onehotCoding)))  
print("Log loss (test) on the VotingClassifier :", log_loss(test_y, vclf.predict_proba(test_x_onehotCoding)))  
print("Number of missclassified point :", np.count_nonzero((vclf.predict(test_x_onehotCoding)- test_y))/test_y.shape[0])  
plot_confusion_matrix(test_y=test_y, predict_y=vclf.predict(test_x_onehotCoding))
```




Log loss (train) on the VotingClassifier : 0.8295219395267811

Log loss (CV) on the VotingClassifier : 1.2063955879095085

Log loss (test) on the VotingClassifier : 1.224902534260774

Number of missclassified point : 0.37293233082706767


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



# Logistic Regression with feature engineering

```

In [93]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----X-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])      Fit the calibrated model
# get_params([deep])      Get parameters for this estimator.
# predict(X)      Predict the target of new samples.
# predict_proba(X)      Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCodingFE, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCodingFE, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCodingFE)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):

```

```
ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty=
'12', loss='log', random_state=42)
clf.fit(train_x_onehotCodingFE, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCodingFE, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCodingFE)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss i
s:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCodingFE)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCodingFE)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss i
s:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```

for alpha = 1e-06
Log Loss : 1.051003134377265
for alpha = 1e-05
Log Loss : 1.0497192969426299
for alpha = 0.0001
Log Loss : 1.0177458951615523
for alpha = 0.001
Log Loss : 0.9901953437928124
for alpha = 0.01
Log Loss : 1.0748573283596188
for alpha = 0.1
Log Loss : 1.3966976110157225
for alpha = 1
Log Loss : 1.8240950106014056
for alpha = 10
Log Loss : 1.8594408627142196
for alpha = 100
Log Loss : 1.8618215887417076

```



For values of best alpha = 0.001 The train log loss is: 0.5233614846923389  
 For values of best alpha = 0.001 The cross validation log loss is: 0.9901953437928124  
 For values of best alpha = 0.001 The test log loss is: 0.9919153581128131

### Testing model with best hyper parameters

```

In [94]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])    Fit linear model with Stochastic Gradient Descent.
# predict(X)    Predict class labels for samples in X.


#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCodingFE, train_y, test_x_onehotCodingFE, test_y, clf)

```


Log loss : 0.9919153581128131

Number of mis-classified points : 0.33383458646616543


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----

**Correctly Classified point**

```
In [95]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty=
'12', loss='log', random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.1672 0.1572 0.0312 0.1138 0.1086 0.1003 0.2854 0.0065 0.0296]]

Actual Class : 7

-----  
Out of the top 500 features 0 are present in query point

### ***Incorrectly Classified point***

```
In [96]: test_point_index = 50
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_
onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_d
f['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.108 0.2367 0.0246 0.1145 0.0527 0.1108 0.3173 0.0062 0.0291]]

Actual Class : 7

-----  
Out of the top 500 features 0 are present in query point

## **Conclusion**

**ngram\_range(1,3), min\_df=3, max\_features=1000 used and below is the output**



```

In [103]: result = pd.DataFrame(columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"])
result = result.append(pd.DataFrame(["Naive Bayes", 0.693, 1.150, 1.247, "36.09%", "Best Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["KNN", 0.600, 0.990, 1.101, "33.08%", "Best Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["LR - Class Balancing", 0.648, 1.136, 1.172, "38.72%", "Best Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["LR -Without Class Balancing", 0.557, 1.165, 1.181, "40.03%", "Over Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["LR -SVM", 0.774, 1.106, 1.180, "33.08%", "Best Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["RF -With OHE", 0.862, 1.186, 1.208, "40.22%", "Over Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["RF -With RC", 0.066, 1.295, 1.339, "46.80%", "Over Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["Stacking Models", 0.686, 1.102, 1.254, "42.25%", "Over Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["Max Voting Classifier", 0.914, 1.160, 1.235, "40.30%", "Over Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame(["LR with FE", 0.640, 0.975, 1.005, "36.39%", "Best Fit"], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result.reset_index(drop = True, inplace = True)
result

```

Out[103]:

	Model	Train Log-loss	CV Log-loss	Test Log-loss	No of Missclassified point	Remarks
0	Naive Bayes	0.693	1.150	1.247	36.09%	Best Fit
1	KNN	0.600	0.990	1.101	33.08%	Best Fit
2	LR - Class Balancing	0.648	1.136	1.172	38.72%	Best Fit
3	LR -Without Class Balancing	0.557	1.165	1.181	40.03%	Over Fit
4	LR -SVM	0.774	1.106	1.180	33.08%	Best Fit
5	RF -With OHE	0.862	1.186	1.208	40.22%	Over Fit
6	RF -With RC	0.066	1.295	1.339	46.80%	Over Fit
7	Stacking Models	0.686	1.102	1.254	42.25%	Over Fit
8	Max Voting Classifier	0.914	1.160	1.235	40.30%	Over Fit
9	LR with FE	0.640	0.975	1.005	36.39%	Best Fit

**In order to reduce test log loss to below 1, made changes as below**

ngram\_range(1,4), min\_df=5, max\_features=3000 we got minimum test log-loss is 0.991 and minimum misclassified points as 33.38% for LR with FE

```

In [104]: result = pd.DataFrame(columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"])
result = result.append(pd.DataFrame([["Naive Bayes", 0.582, 1.242, 1.264, "40.03%", "Over Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["KNN", 0.869, 1.081, 1.088, "37.03%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["LR - Class Balancing", 0.554, 1.240, 1.242, "38.90%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["LR -Without Class Balancing", 0.606, 1.222, 1.235, "40.97%", "Over Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["LR -SVM", 0.417, 1.008, 1.086, "32.14%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["RF -With OHE", 0.529, 1.185, 1.182, "38.90%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["RF -With RC", 0.069, 1.296, 1.292, "41.91%", "Over Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["Stacking Models", 0.459, 1.243, 1.275, "38.49%", "Over Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["Max Voting Classifier", 0.082, 1.206, 1.224, "37.29%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result = result.append(pd.DataFrame([["LR with FE", 0.523, 0.990, 0.991, "33.38%", "Best Fit"]], columns = ["Model", "Train Log-loss", "CV Log-loss", "Test Log-loss", "No of Missclassified point", "Remarks"]))
result.reset_index(drop = True, inplace = True)
result

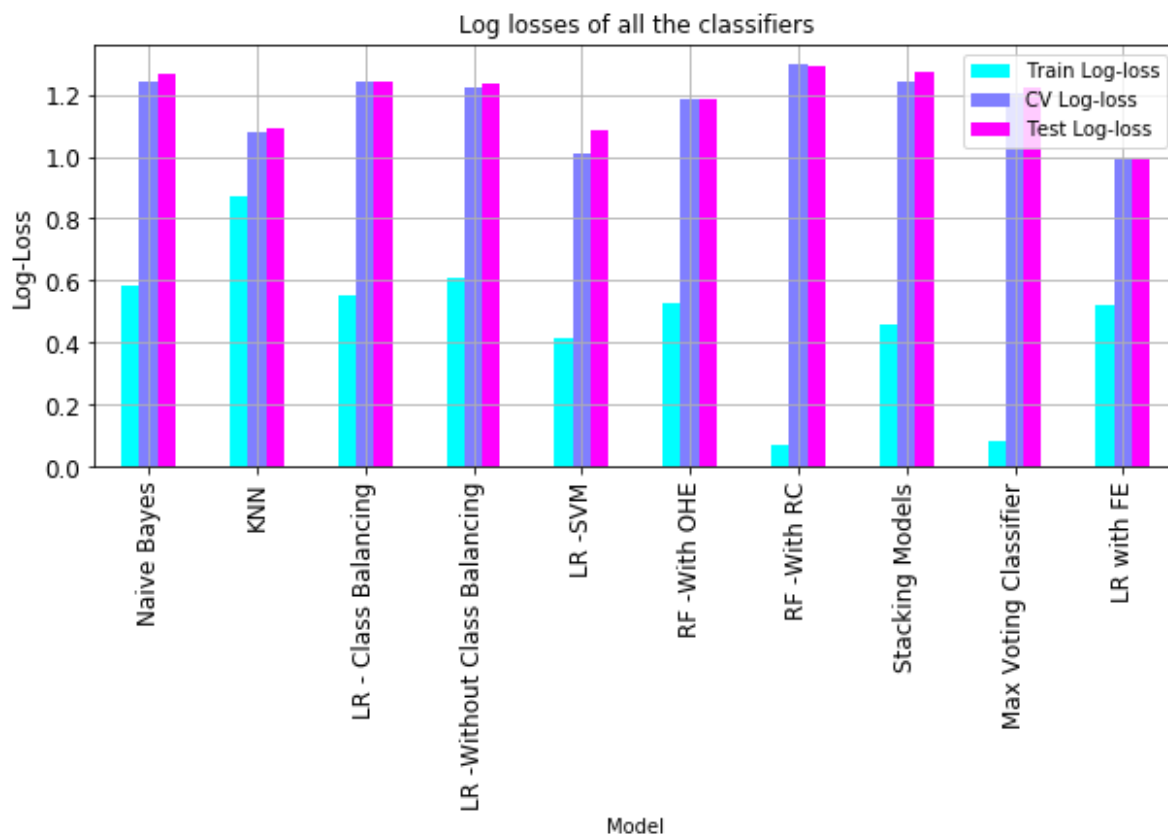
```

Out[104]:

	Model	Train Log-loss	CV Log-loss	Test Log-loss	No of Missclassified point	Remarks
0	Naive Bayes	0.582	1.242	1.264	40.03%	Over Fit
1	KNN	0.869	1.081	1.088	37.03%	Best Fit
2	LR - Class Balancing	0.554	1.240	1.242	38.90%	Best Fit
3	LR -Without Class Balancing	0.606	1.222	1.235	40.97%	Over Fit
4	LR -SVM	0.417	1.008	1.086	32.14%	Best Fit
5	RF -With OHE	0.529	1.185	1.182	38.90%	Best Fit
6	RF -With RC	0.069	1.296	1.292	41.91%	Over Fit
7	Stacking Models	0.459	1.243	1.275	38.49%	Over Fit
8	Max Voting Classifier	0.082	1.206	1.224	37.29%	Best Fit
9	LR with FE	0.523	0.990	0.991	33.38%	Best Fit

```
In [101]: bar_result = result.drop(["No of Missclassified point", "Remarks"], axis = 1)
```

```
In [102]: bar_result.plot(x = "Model", kind = "bar", figsize = (10, 4), grid = True, fontsize = 12, colormap="cool")
plt.title("Log losses of all the classifiers", fontsize = 12)
plt.ylabel("Log-Loss", fontsize = 12)
plt.show()
```



## Summary

1. Applied all the models with tf-idf feature
2. Maximum 1000 words used
3. NB - OHE is applied
4. KNN - Response coding is applied
5. LR with count vectorizer and unigram and bigram and OHE is applied, with class balancing
6. LR with count vectorizer and unigram and bigram and OHE is applied, without class balancing
7. Lr SVM with class balancing with OHE is applied
8. RF with OHE
9. RF with response coding
10. stacking the models with OHE
11. Maximum voting classifier with OHE
12. LR with FE